
MDAgents: An Adaptive Collaboration of LLMs for Medical Decision-Making

Yubin Kim¹ Chanwoo Park¹ Hyewon Jeong^{1‡} Yik Siu Chan¹

Xuhai Xu¹ Daniel McDuff² Hyeonhoon Lee³

Marzyeh Ghassemi¹ Cynthia Breazeal¹ Hae Won Park¹

¹Massachusetts Institute of Technology

²Google Research

³Seoul National University Hospital

{ybkim95, cpark97, hyewonj, yiksiuc, xoxu, mghassem, cynthiab, haewon}@mit.edu

dmcduff@google.com

hhoon@snu.ac.kr

Abstract

Foundation models are becoming valuable tools in medicine. Yet despite their promise, the best way to leverage Large Language Models (LLMs) in complex medical tasks remains an open question. We introduce a novel multi-actor framework, named Medical Decision-making Agents (MDAgents) that helps to address this gap by automatically assigning a collaboration structure to a team of LLMs. The assigned solo or group collaboration structure is tailored to the medical task at hand, a simple emulation inspired by the way real-world medical decision-making processes are adapted to tasks of different complexities. We evaluate our framework and baseline methods using state-of-the-art LLMs across a suite of real-world medical knowledge and clinical diagnosis benchmarks. MDAgents achieved the **best performance in seven out of ten** benchmarks on tasks requiring an understanding of medical knowledge and multi-modal reasoning, showing a significant **improvement of up to 6.5%** ($p < 0.05$) compared to previous methods' best performances. Ablation studies reveal that MDAgents effectively determines medical complexity to optimize for efficiency and accuracy across diverse medical tasks. Notably, the combination of moderator review and external medical knowledge in group collaboration resulted in an average accuracy improvement of 11.8%. Our code can be found at <https://anonymous.4open.science/r/MDAgents>.

1 Introduction

Medical Decision-Making (MDM) is a multifaceted and intricate process in which clinicians collaboratively navigate diverse sources of information to reach a precise and specific conclusion [90]. For instance, a primary care physician (PCP) may refer a patient to a specialist when faced with a complex case, or a patient visiting the emergency department or urgent care might be triaged and then directed to a specialist for further evaluation [5, 49]. MDM involves interpreting complex and multi-modal data, such as imaging, electronic health records (EHR), physiological signals, and genetic information, while rapidly integrating new medical research into clinical practice [62, 72]. Recently, Large Language Models (LLMs) have shown potential for AI support in MDM [22, 58, 66]. Their ability to process and synthesize large volumes of medical literature [68] and clinical information [1], as well as support probabilistic [87] and causal [37] reasoning, makes LLMs promising tools. However, there is no silver bullet, and medical applications especially require careful design.

[‡]Hyewon Jeong received her MD degree from Yonsei University College of Medicine, South Korea.

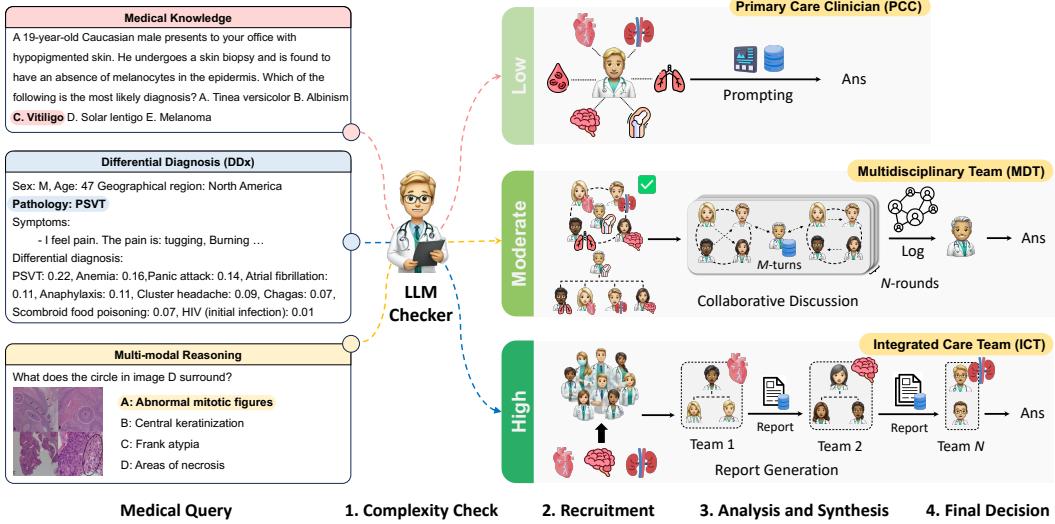


Figure 1: **Medical Decision-Making Agents (MDAgents) framework.** Given a medical query from different medical datasets, the framework performs 1) medical complexity check, 2) recruitment, 3) analysis and synthesis, and 4) decision-making steps.

While decision-making tools including multi-agent LLMs [11, 81] have shown promise in non-medical domains [31, 32, 41, 43, 56, 59], their evaluation in health applications has been limited. To date, their “generalist” design has not effectively integrated the real-world systematic MDM process [51] which requires an adaptive, collaborative, and tiered approach. Clinicians consider the current and past history of the patient, available evidence from medical literature, and their domain expertise and experience [20] for MDM. One example of MDM is to triage patients in urgent care units based on the severity and complexity of their medical conditions [12, 26, 82]. Patients with pathognomonic, single uncomplicated acute conditions, or stable chronic conditions that their PCP could manage [80] could be low complexity cases. On the other hand, patients with complicated injuries, chronic conditions with side effects, or superimposed diseases who often require consultation with specialty physicians [27, 55] are considered high complexity cases².

Inspired by the way that clinicians make decisions in practice, we propose **Medical Decision-making Agents (MDAgents)**, an adaptive medical decision-making framework that leverages LLMs to emulate the hierarchical diagnosis procedures ranging from individual clinicians to collaborative clinician teams (Figure 1). MDAGents work in three steps: 1) Medical complexity check; 2) Recruitment based on medical complexity; 3) Analysis and synthesis and 4) Final decision-making to return the answer. Our contributions are threefold:

1. We introduce MDAGents, the first adaptive decision-making framework for LLMs that mirrors real-world MDM processes via dynamic collaboration among AI agents based on task complexity.
2. MDAGents demonstrate superior performance in accuracy over previous solo and group methods on 7 out of 10 medical benchmarks, and we show an effective trade-off between performance and efficiency (i.e., the number of API calls) by varying the number of agents.
3. We provide rigorous testing under various hyperparameters (e.g., temperatures), demonstrating better robustness of MDAGents compared to solo and group methods. Furthermore, our ablations evidence MDAGents’ ability to find the appropriate complexity level for each MDM instance.

2 Related Works

Language Models in Medical Decision-Making LLMs have shown promise in a range of applications within the medical field [14, 38, 57, 69, 70, 89]. They can answer questions from medical exams [40, 47], perform biomedical research [35], and clinical diagnosis [50, 61]. Medical LLMs are also evaluated on generative tasks, including creating medical reports [73], describing medical images [75], constructing differentials [48], performing diagnostic dialogue with patients [71], and generating psychiatric evaluations of interviews [24]. To advance the capabilities of medical LLMs, two main approaches have been explored: (1) training with domain-specific data [28], and (2) applying inference-time techniques such as prompt engineering [61] and Retrieval Augmented Generation

²Detailed examples of low-, moderate- and high-complexity cases are provided in Appendix D.1.

Table 1: Comparison between our framework and previous methods (Solo and Group). Among these works, MDAgents is the only one to perform all key dimensions of LLM decision-making.

Method	MDAgents (Ours)	Single	Voting [77]	Debate [17]	MedAgents [66]	ReConcile [10]
Interaction Type						
Multiple Roles	✓	✗	✓	✓	✓	✓
Early Stopping	✓	✗	✓	✓	✓	✗
Refinement	✓	✗	✗	✓	✓	✗
Complexity Check	✓	✗	✗	✗	✗	✗
Multi-party Chat	✓	✗	✗	✓	✗	✗
Conversation Pattern	Flexible	Static	Static	Static	Static	Static

(RAG) [85]. While initial research has been concentrated on pre-training and fine-tuning with medical knowledge, the rise of large general-purpose LLMs has enabled training-free methods where models leverage their latent medical knowledge. For example, GPT-4 [54], with richer prompt crafting, surpasses the passing score on USMLE by over 20 points and with prompt tuning can outperform fine-tuned models including Med-PaLM [52, 53]. The promise of general-purpose models has thus inspired various techniques such as Medprompt and ensemble refinement to improve LLM reasoning [61], as well as RAG tools that use external resources to improve the factuality and completeness of LLM responses [36, 85]. Our approach leverages these techniques and the capabilities of general-purpose models while acknowledging that a solitary LLM may not fully encapsulate the collaborative and multidisciplinary nature of real-world MDM. We thus emphasize joining multiple expert LLMs for effective collaboration in order to solve complicated medical tasks with greater accuracy.

Multi-Agent Collaboration An array of studies have explored effective collaboration frameworks between multiple LLM agents [44, 81] to enhance capability above and beyond an individual LLM [74]. A common framework is role-playing, where each agent adopts a specific role (e.g. an Assistant Agent or a Manager Agent), a task is then broken down into sub-steps and solved collaboratively [44, 81]. While role-playing focuses on collaboration and multi-step problem-solving [83], another framework, “multi-agent debate”, prompts each agent to solve the task independently [17]. Then, they reason through other agents’ answers to converge on a shared response, this approach can improve the factuality, mathematical ability and reasoning capabilities of the multi-agent solution [17, 45]. Similar frameworks include voting [77], multi-disciplinary collaboration [66], group discussions (ReConcile [10]), and negotiating [23]. Table 1 compares existing setups across key dimensions in multi-agent interaction. Although these frameworks have shown improvement in the respective tasks, they rely on a pre-determined number of agents and interaction settings. When applied on a wider variety of tasks , this static architecture may lead to suboptimal multi-agent configurations, negatively impacting performance [46]. Furthermore, multi-agent approaches run the risk of being computationally inefficient or expensive to employ and need to justify these costs with noticeable performance gains [17]. Given that different models and frameworks could generalize better to different tasks [86], we propose a framework that dynamically assigns the optimal collaboration strategy at inference time based on the complexity of the query. We apply our strategy to MDM, a task that requires teamwork and should benefit from multi-agent collaboration [66].

3 MDAgents: Medical Decision-making Agents

The design of MDAgents (Figures 1 and 2) incorporates four stages: **1) Medical Complexity Check** - The system evaluates the medical query, categorizing it as *low*, *moderate*, or *high* complexity based on clinical decision-making techniques [6, 7, 21, 65, 79]. **2) Expert Recruitment** - Based on complexity, the framework activates a single Primary Care Clinician (PCC) for low complexity issues, or a Multi-disciplinary Team (MDT) or Integrated Care Team (ICT) for moderate or high complexities [7, 18, 21, 33, 42, 65]. **3) Analysis and Synthesis** - Solo queries use prompting techniques like Chain-of-Thought (CoT) and Self-Consistency (SC). MDTs involve multiple LLM agents forming a consensus, while ICTs synthesize information for the most complex cases. **4) Decision-making** - The final stage synthesizes all inputs to provide a well-informed answer to the medical query.

3.1 Agent Roles

Moderator. The moderator agent functions as a general practitioner (GP) or emergency department doctor who first triages the medical query. This agent assesses the complexity of the problem and determines whether it should be handled by a single agent, a MDT, or an ICT. The moderator

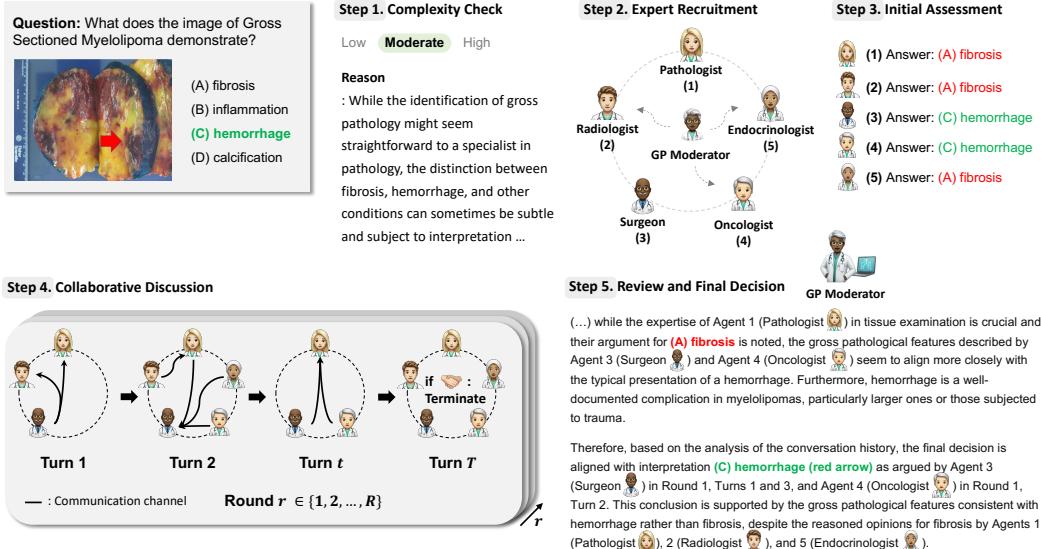


Figure 2: Illustrative example of MDAgents in a *moderate* complexity case from the PMC-VQA dataset. More detailed case studies can be found in Figure 11 and 12 in the Appendix.

ensures the appropriate pathway be selected based on the query’s complexity and oversees the entire decision-making process.

Recruiter. The recruiter agent is responsible for assembling the appropriate team of specialist agents based on the complexity assessment of the moderator. The recruiter may assign a single PCP agent for low-complexity cases, while MDT or ICT with relevant expertise will be formed for moderate and high-complexity cases.

General Doctor/Specialist. These agents are domain-specific or general physicians recruited by the recruiter agent. Depending on the complexity of the case, they may work independently or as part of a team. General physicians handle less complex, routine cases, whereas specialists are recruited for their specific expertise in more complex scenarios. These agents engage in the collaborative decision-making process, contributing their specialized knowledge to reach a consensus or provide detailed reports for high-complexity cases.

3.2 Medical Complexity Classification (Line 1 of Algorithm 1, Appendix C.2)

The first step in the MDAgents framework is to determine the complexity of a given medical query q by the *moderator LLM* which functions as a *generalist practitioner* (GP). The moderator aims to act as a classifier to return the complexity level of the given medical query, it is provided with the information on how medical complexity should be defined and is instructed to classify the query into one of three different complexity levels:

1. **Low** - Simple, well-defined medical issues that can be resolved by a single PCP agent. These typically include common, acute illnesses or stable chronic conditions where the medical needs are predictable and require minimal interdisciplinary coordination.
2. **Moderate** - The medical issues involve multiple interacting factors, necessitating a collaborative approach among an MDT. These scenarios require integration of diverse medical knowledge areas and coordination between specialists through consultation to develop effective care strategies.
3. **High** - Highly complex medical scenarios that demand extensive coordination and combined expertise from an ICT. These cases often involve multiple chronic conditions, complicated surgical or trauma cases, and complex decision-making that integrates various healthcare departments and multiple specialists.

3.3 Expert Recruitment (Line 3, 7, 17 of Algorithm 1)

Given a medical query, the goal of the *recruiter LLM* is to enlist domain experts as individuals, in groups, or as multiple teams, based on the complexity levels determined by the *moderator LLM*. Specifically, we assign medical expertise and roles to multiple LLMs, instructing them to either act independently as solo medical agents or collaborate with other medical experts in a team. In Figure 9 in the Appendix, we also provide frequently recruited agents for each benchmark as a reference.

3.4 Medical Collaboration and Refinement

The initial assessment protocol of our decision-making framework categorizes query complexity into *low*, *moderate*, and *high*. This categorization is grounded in established medical constructs such as

acuity [25] for straightforward cases, comorbidity [63] and case management complexity [13] for intermediate and multi-disciplinary care requirements, and severity of illness [16] for high complexity cases requiring comprehensive management. We outline the specific refinement approach:

Low - Straightforward cases (Line 2-4 of Algorithm 1). For queries classified under Low complexity, characterized by straightforward clinical decision pathways, a single PCP agent (Figure 10 (a)) is deployed by the definition in Section 3.2. The domain expert who is recruited by *recruiter* LLM, applies few-shot prompting to the problem. The output answer, denoted as ans , is directly obtained from the agent's response to Q without the need for iterative refinement, formalized as $ans = Agent(Q)$, with $Agent$ representing the engaged PCP agent.

Moderate - Intermediate complexity cases (Line 6-14 of Algorithm 1). In addressing more complex queries, the utilization of an MDT (Figure 10 (b) and (c)) approach has been increasingly recognized for its effectiveness in producing comprehensive and nuanced solutions [42]. The MDT framework leverages the collective expertise of professionals from diverse disciplines, facilitating a holistic examination of the query at hand. This collaborative method is particularly advantageous in scenarios where the complexity of a problem transcends the scope of a single domain, necessitating a fusion of insights from various specialties [7, 65]. The MDT approach not only enhances decision-making quality through the integration of multidimensional perspectives but also significantly improves the adaptability and efficiency of the problem-solving process [21].

Building upon this foundation, our framework specifically addresses queries of moderate complexity through a structured, multi-tiered collaborative approach. An MDT recruited by *recruiter* LLM (see Figure 10 in Appendix) starts an iterative discussion process aimed at reaching a consensus with at most R rounds (Line 10-12). For every round, consensus within the MDT is determined by parsing and comparing their opinions. In the event of a disagreement, the moderator agent, consistent with the one described in Section 3.2 reviews the MDT's discourse and formulates feedback for each agent.

High - Complex care cases (Line 17-24 of Algorithm 1). In contrast to the MDT approach, the ICT (Figure 10 (d)) paradigm is essential for addressing the highest tier of query complexity in healthcare. This structured progression through the ICT ensures a depth of analysis that is specialized and focused at each stage of the decision-making process. Beginning with the Initial Assessment Team, moving through various diagnostic teams, and culminating with the Final Review & Decision Team, our ICT model aligns specialist insights into a cohesive narrative that informs the ultimate decision (Appendix Algorithms 1 Lines 19-21). This phased approach, supported by evidence from recent healthcare studies, has been shown to enhance the precision of clinical decision-making, as each team builds upon the foundation laid by the previous, ensuring a meticulous and refined examination of complex medical cases [33]. The resultant reports are not only reflective of comprehensive medical evaluations but also of a systematic and layered analysis that is critical in the management of intricate health scenarios [18].

3.5 Decision-making

In the final stage of our framework, the decision-maker LLM agent synthesizes the diverse inputs generated throughout the decision-making process to arrive at a well-informed final answer to the medical query q . This synthesis involves several components depending on the complexity level of the query:

1. *Low*: Directly utilizes the initial response from the primary decision-making agent.
2. *Moderate*: Incorporates the conversation history (*Interaction*) between the recruited agents to understand the nuances and disagreements in their responses.
3. *High*: Considers detailed reports (*Reports*) generated by the agents, which include comprehensive analyses and justifications for their diagnostic suggestions.

The decision-making process is formulated as $ans = Agent(\cdot)$ where the final answer, ans is determined by integrating the outputs from analysis and synthesis step based on its medical complexities. This integration employs ensemble techniques such as temperature ensembles to ensure the decision is robust and reflects a consensus among the models when applicable (see Appendix C.2 for details).

4 Experiments and Results

In this section, we evaluate our framework and baseline methods across different medical benchmarks in **Solo**, **Group**, and **Aaptive** settings. Our experiments and ablation studies highlight the framework's performance, demonstrating robustness and efficiency by modulating agent numbers and temperatures.

4.1 Setup Results also show a beneficial convergence of agent opinions in collaborative settings.

To verify the effectiveness of our framework, we conduct comprehensive experiments with baseline methods on ten datasets including MedQA, PubMedQA, DDXPlus, SymCat, JAMA, MedBullets, Path-VQA, PMC-VQA, MIMIC-CXR and MedVidQA. A detailed explanation and statistics for each

Table 2: Accuracy (%) on Medical benchmarks with **Solo**/**Group**/**Adaptive** setting. **Bold** represents the best and Underlined represents the second best performance for each benchmark and model. All benchmarks except for MedVidQA were evaluated with GPT-4(V) and MedVidQA was evaluated with Gemini-Pro(Vision). Full experimental results with other models are listed in Table 5-7.

		Medical Knowledge Retrieval Datasets				
Category	Method	MedQA	PubMedQA	Path-VQA	PMC-VQA	MedVidQA
Single-agent	Zero-shot	75.0 \pm 1.3	61.5 \pm 2.2	57.9 \pm 1.6	49.0 \pm 3.7	37.9 \pm 8.4
	Few-shot	72.9 \pm 11.4	63.1 \pm 11.7	57.5 \pm 4.5	52.2 \pm 2.0	47.1 \pm 8.6
	+ CoT [78]	82.5 \pm 4.9	57.6 \pm 9.2	58.6 \pm 3.1	51.3 \pm 1.5	48.6 \pm 5.5
	+ CoT-SC [77]	83.9 \pm 2.7	58.7 \pm 5.0	61.2 \pm 2.1	50.5 \pm 5.2	49.2 \pm 8.2
	ER [61]	81.9 \pm 2.1	56.0 \pm 7.0	61.4 \pm 4.1	52.7 \pm 2.9	48.5 \pm 4.1
Multi-agent (Single-model)	Majority Voting	80.6 \pm 2.9	72.2 \pm 6.9	56.9 \pm 19.7	36.8 \pm 6.7	50.8 \pm 7.4
	Weighted Voting	78.8 \pm 1.1	72.2 \pm 6.9	62.1 \pm 13.9	25.4 \pm 9.0	57.8 \pm 2.1
	Borda Count	70.3 \pm 8.5	66.9 \pm 3.0	61.9 \pm 8.1	27.9 \pm 5.3	54.5 \pm 4.7
	MedAgents [66]	79.1 \pm 7.4	69.7 \pm 4.7	45.4 \pm 8.1	39.6 \pm 3.0	51.6 \pm 4.8
	Meta-Prompting [64]	80.6 \pm 1.2	73.3 \pm 2.3	55.3 \pm 2.3	42.6 \pm 4.2	-
Multi-agent (Multi-model)	Reconcile [10]	81.3 \pm 3.0	79.7 \pm 3.2	57.5 \pm 3.3	31.4 \pm 1.2	-
	AutoGen [81]	60.6 \pm 5.0	77.3 \pm 2.3	43.0 \pm 8.9	37.3 \pm 6.1	-
	DyLAN [46]	64.2 \pm 2.3	73.6 \pm 4.2	41.3 \pm 1.2	34.0 \pm 3.5	-
Adaptive	MDAgents (Ours)	88.7 \pm 4.0	75.0 \pm 1.0	65.3 \pm 3.9	56.4 \pm 4.5	56.2 \pm 6.7
Clinical Reasoning and Diagnostic Datasets						
Category	Method	DDXPlus	SymCat	JAMA	MedBullets	MIMIC-CXR
Single-agent	Zero-shot	70.3 \pm 2.0	88.7 \pm 2.3	62.0 \pm 2.0	67.0 \pm 1.4	40.0 \pm 5.3
	Few-shot	69.4 \pm 1.0	86.7 \pm 3.1	69.0 \pm 4.2	72.0 \pm 2.8	35.3 \pm 5.0
	+ CoT [78]	72.7 \pm 7.7	78.0 \pm 2.0	66.0 \pm 5.7	70.0 \pm 0.0	36.2 \pm 5.2
	+ CoT-SC [77]	52.1 \pm 6.4	83.3 \pm 3.1	68.0 \pm 2.8	76.0 \pm 2.8	51.7 \pm 4.0
	ER [61]	61.3 \pm 2.4	82.7 \pm 2.3	71.0 \pm 1.4	76.0 \pm 5.7	50.0 \pm 0.0
Multi-agent (Single-model)	Medprompt [53]	59.5 \pm 17.7	87.3 \pm 1.2	70.7 \pm 4.3	71.0 \pm 1.4	53.4 \pm 4.3
	Majority Voting	67.8 \pm 4.9	91.9 \pm 2.2	70.0 \pm 5.7	70.0 \pm 0.0	49.5 \pm 10.7
	Weighted Voting	65.9 \pm 3.3	90.5 \pm 2.9	66.1 \pm 4.1	66.0 \pm 5.7	53.5 \pm 2.2
	Borda Count	67.1 \pm 6.7	78.0 \pm 11.8	61.0 \pm 5.6	66.0 \pm 2.8	45.3 \pm 6.8
	MedAgents [66]	62.8 \pm 5.6	90.0 \pm 0.0	66.0 \pm 5.7	77.0 \pm 1.4	43.3 \pm 7.0
Multi-agent (Multi-model)	Meta-Prompting [64]	52.6 \pm 6.1	77.3 \pm 2.3	64.7 \pm 3.1	49.3 \pm 1.2	42.0 \pm 4.0
	Reconcile [10]	68.4 \pm 7.4	90.6 \pm 2.5	60.7 \pm 5.7	59.5 \pm 8.7	33.3 \pm 3.4
	AutoGen [81]	67.3 \pm 11.8	73.3 \pm 3.1	64.6 \pm 1.2	55.3 \pm 3.1	43.3 \pm 4.2
	DyLAN [46]	56.4 \pm 2.9	75.3 \pm 4.6	60.1 \pm 3.1	57.3 \pm 6.1	38.7 \pm 1.2
Adaptive	MDAgents (Ours)	77.9 \pm 2.1	93.1 \pm 1.0	70.9 \pm 0.3	80.8 \pm 1.7	55.9 \pm 9.1

* CoT: Chain-of-Thought, SC: Self-Consistency, ER: Ensemble Refinement

* text-only, image+text, video+text

* All experiments were tested with 3 random seeds

dataset are deferred to Appendix A and Figure 8. We use 50 samples per dataset for testing, and the inference time for each complexity was: *low* - 14.7s, *moderate* - 95.5s, and *high* - 226s in average. For all the quantitative experiments in this section, we compare three settings: (1) **Solo**: Using a single LLM agent in the decision-making state. (2) **Group**: Implementing multi-agents to collaborate during the decision-making process. (3) **Adaptive**: Our proposed method MDAgents, adaptively constructs the inference structure from PCP to MDT and ICT.

Medical Question Answering With MedQA [34], PubMedQA [35], MedBullets [9], and JAMA [9], we focus on question answering through text, involving both literature-based and conceptual medical knowledge questions. Specifically, PubMedQA tasks models to answer questions using abstracts from PubMed, requiring synthesis of biomedical information. MedQA tests the model's ability to understand and respond to multiple-choice questions derived from medical educational materials and examinations. MedBullets provides USMLE Step 2/3 type questions that demand the application of medical knowledge and clinical reasoning. JAMA Clinical Challenge presents challenging real-world clinical cases with diagnosis or treatment decision-making questions, testing the model's clinical reasoning (Figure 8 in Appendix shows complexity distribution for each dataset)

Diagnostic Reasoning DDXPlus [67] and SymCat [2] involve clinical vignettes that require differential diagnosis, closely mimicking the diagnostic process of physicians. These tasks test the model's ability to reason through symptoms and clinical data to suggest possible medical conditions, evaluating the AI's diagnostic reasoning abilities similar to a clinical setting. SymCat [2] uses synthetic patient records constructed from a public disease-symptom data source and is enhanced with additional contextual information through the NLICE method.

Medical Visual Interpretation Path-VQA [30], PMC-VQA [88], MedVidQA [29], and MIMIC-CXR [3] challenge models to interpret medical images and videos, requiring integration of visual data with clinical knowledge. PathVQA focuses on answering questions based on pathology images, testing AI's capability to interpret complex visual information from medical images. PMC-VQA

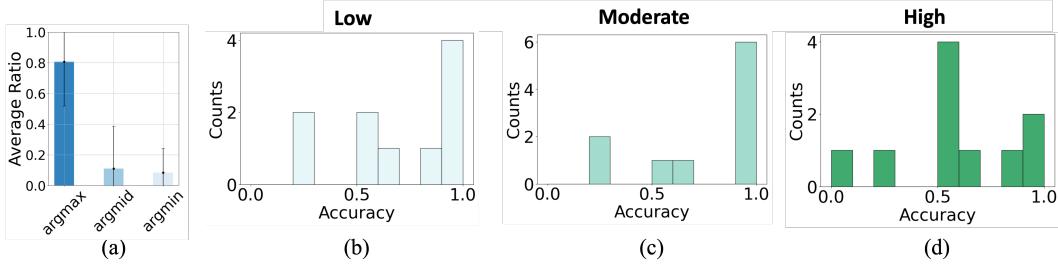


Figure 3: Experiment with MedQA dataset with $N=10$ randomly sampled questions. (a) The average plot for the probability of moderator provides the complexity level with the best / middle / worst choice for each problem. (b-d) The distribution of accuracy level of low / moderate / high solver. We solved each question with a specific complexity level 10 times.

evaluates AI’s proficiency in deriving answers from both text and images found in scientific publications. MedVidQA extends to video-based content, where AI models need to process information from medical procedure videos. MIMIC-CXR-VQA specifically targets chest radiographs, utilizing a diverse and large-scale dataset designed for visual question-answering tasks in the medical domain.

Baseline Methods

- **Solo:** The baseline methods considered for the Solo setting include the following: Zero-shot [39] directly incorporates a prompt to facilitate inference, while Few-shot [8] involves a small number of examples. Few-shot CoT [78] integrates rationales before deducing the answer. Few-shot CoT-SC [76] builds upon Few-shot CoT by sampling multiple chains to yield the majority answer. Ensemble Refinement (ER) [61] is a prompting strategy that conditions model responses on multiple reasoning paths to bolster the reasoning capabilities of LLMs. Medprompt [53] is a composition of several prompting strategies that enhances the performance of LLMs and achieves state-of-the-art results on multiple benchmark datasets, including medical and non-medical domains.
- **Group:** We tested five group decision-making methods: Voting [77], MedAgents [66], Reconcile [10], AutoGen [81], and DyLAN [46]. Autogen was based on four agents, with one User, one Clinician, one Medical Expert, and one Moderator, with one response per agent [81]. DyLAN setup followed the base implementations of four agents with no specific roles and four maximum rounds of interaction [46]. While the methods support multiple models, GPT-4 was used for all agents.

4.2 Results

In Table 2, we report the classification accuracy on MedQA, PubMedQA, DDXPlus, SymCat, JAMA, MedBullets, Path-VQA, PMC-VQA and MedVidQA dataset. We compare our method (Adaptive) with several baselines in both Solo and Group settings.

Adaptive method outperforms Solo and Group settings. As depicted in Figure 4 and Table 2, MDAGents significantly outperforms ($p < 0.05$) both Solo and Group setting methods, showing best performance in **7 out of 10** medical benchmarks tested. This reveals the effectiveness of adaptive strategies integrated within our system, particularly when navigating through the text-only (e.g., DDXPlus where it outperformed the best performance of single-agent by 7.2% and multi-agent by 9.5%) and text-image datasets (e.g., Path-VQA, PMC-VQA and MIMIC-CXR). Our approach not only comprehends textual information with high precision but also adeptly synthesizes visual data, a pivotal capability in medical diagnostic evaluations.

Why Adaptive Decision-making Framework Works Well We hypothesize that an LLM, functioning as a classifier or a generalist doctor, will select the optimal complexity level for each MDM problem. This hypothesis is supported by Figure 3, which illustrates that the model appropriately matches the complexity levels; low, moderate, and high of the given problem.

Formally, for any given problem P , we denote the likelihood that the correct answer can be solved by a specific complexity level as $p_{\text{complexity-level}}(P)$, where $\text{complexity-level} \in \{\text{low}, \text{moderate}, \text{high}\}$. The classifier LLM selects the complexity level according to the highest, middle, and lowest probabilities, referred to as arg max , arg mid , and arg min respectively, with probabilities a , b , and c . Therefore, the accuracy of our system for problem P can be described by the equation $a \cdot p_{\text{argmax}}(P) + b \cdot p_{\text{argmid}}(P) + c \cdot p_{\text{argmin}}(P)$, and also the overall accuracy would be $\mathbb{E}_P[a \cdot p_{\text{argmax}}(P) + b \cdot p_{\text{argmid}}(P) + c \cdot p_{\text{argmin}}(P)]$. The calculated values of a , b , c are $a = 0.81 \pm 0.29$, $b = 0.11 \pm 0.28$, and $c = 0.08 \pm 0.16$, which indicates that LLM can provide an optimal complexity level with probability at least 80%.

These findings suggest that a classifier LLM can implicitly simulate various complexity levels and optimally adapt to the complexity required for each medical problem, as shown in Figure 3. This ability to adjust complexity dynamically proves to be crucial for applying LLMs effectively in MDM contexts as shown by the competitiveness of our Adaptive approach.

Solo vs. Group Setting in MDM. The experimental results reveal distinct performance patterns between Solo and Group settings across various medical benchmarks. In simpler datasets like MedQA, solo methods, leveraging Few-shot CoT and CoT-SC, achieved up to 83.9% accuracy compared to the group’s best of 81.3%. Conversely, for more complex datasets like SymCat, group settings perform better, with SymCat showing 91.9% accuracy in the group settings versus 88.7% in solo settings. Notably, group settings (e.g. Weighted Voting, Reconcile) performed better in multi-modal datasets such as Path-VQA (*image + text*), MedVidQA (*video + text*), and MIMIC-CXR (*image + text*), highlighting the advantage of collaborative process in complex tasks. This result aligns with findings from [4], which showed that pooled diagnoses from groups of physicians significantly outperformed those of individual physicians, with accuracy increasing as group size increased. Overall, solo settings outperformed group settings in four benchmarks, while group settings outperformed solo in six benchmarks. These results reveals that while solo methods excel in straightforward tasks, group settings provide better performance in complex, multi-faceted tasks requiring diverse expertise.

4.3 Ablation Studies

Impact of Complexity Selection. We evaluate the importance of the complexity assessment and adaptive process through ablation studies (Figure 5). Our adaptive method significantly outperforms static complexity assignments across different modality benchmarks. For *text-only* queries, the Adaptive method achieves an accuracy of 81.2%, significantly higher than *low* (64.2%), *moderate* (71.6%), and *high* (65.8%) settings. Interestingly, 64% of the text-only queries were classified as *high* complexity, indicating that many text-based queries required in-depth analysis with different expertise. In the *image + text* modality, the Adaptive method classified 55% of the queries as *low* complexity, suggesting that the visual information often provides clear and straightforward cues that simplify the decision-making process. Finally, for *video + text* queries, 87% of these queries were classified as *low* complexity, reflecting that the dynamic visual data in conjunction with text can often be straightforwardly interpreted. However, further evaluation on more challenging video medical datasets is needed, as MedVidQA contains relatively less complex medical knowledge.

Impact of Moderator’s Review and RAG Table 3 examines the impact of incorporating external medical knowledge and moderator reviews into the MDAgents framework on accuracy. MedRAG [84] is a systematic toolkit for Retrieval-Augmented Generation (RAG) that leverages various corpora; biomedical, clinical and general medicine, to provide comprehensive knowledge. The baseline accuracy of MDAgents is 71.8%. Integrating MedRAG increases accuracy to 75.2% (up 4.7%), while the moderator’s review alone raises it to 77.6% (up 8.1%). The combined use of both methods achieves the highest accuracy at 80.3% (up 11.8%).

The results indicate that MedRAG and moderator review both enhance performance, with their combined effect being synergistic. This highlights that leveraging recent external knowledge and structured feedback mechanisms is crucial for refining and converging on accurate medical decisions. This improvement underscores the importance of a hybrid strategy, aligning with real-world practices of continuous learning and expert consultation to optimize performance in medical applications.

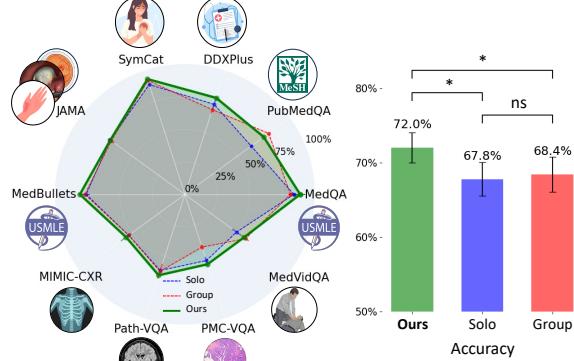


Figure 4: Our method outperforms Solo and Group settings across different medical benchmarks.

Method	Avg. Accuracy (%)
MDAgents (Ours)	71.8
+ MedRAG	75.2 (↑ 4.7 %)
+ Moderator’s Review	77.6 (↑ 8.1 %)
+ Moderator’s Review & MedRAG	80.3 (↑ 11.8 %)

Table 3: Ablations for the impact of moderator’s review and MedRAG. The Accuracy were averaged accuracy across all datasets.

4.4 Impact of Number of Agents in Group Setting.

Our experiment with varying the number of agents in a collaborative Group setting (Appendix Figure 6 (a-b)) shows that a higher number of agents does not lead to better performance. Rather, our Adaptive method achieves optimal performance with fewer agents ($N=3$, peak accuracy of 83.5%) by intelligently calibrating the number of collaborating agents. This not only indicates efficiency in

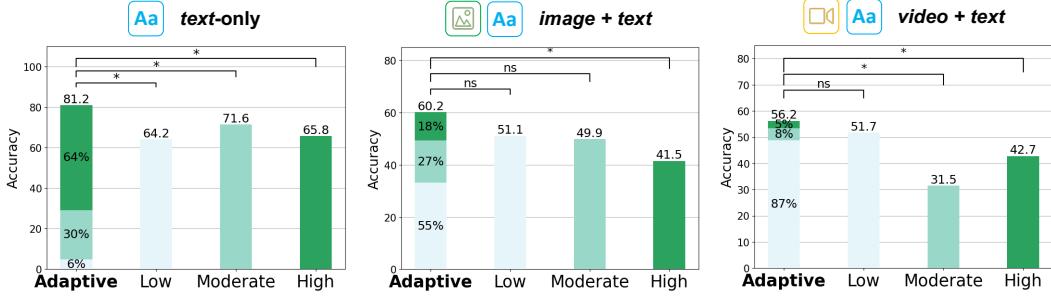


Figure 5: Impact of complexity selection of the query. Accuracy of each ablation on *text-only* (left), *text+image* (center) and *text+video* (right) benchmarks are reported.

decision-making but also computational and economic benefits, considering the reduced number of API calls needed, especially when contrasted with the Solo and Group settings.

With regards to computational efficiency, the Solo setting (5-shot CoT-SC) resulted in a 6.0 and Group setting (MedAgents with $N=5$) resulted in a 20.3 API calls, suggesting a high computational cost without a corresponding increase in accuracy. On the other hand, our Adaptive method exhibits a more economical use of resources, demonstrated by fewer API calls (9.3 with $N=3$) while maintaining high accuracy, a critical factor in deploying scalable and cost-effective medical AI solutions.

4.5 Robustness of MDAGENTS with different parameters.

Our Adaptive approach shows resilience to changes in temperature (Appendix Figure 6 (c), low ($T=0.3$) and high ($T=1.2$)) with performance improving under higher temperatures. This suggests that our model can utilize the creative and diverse outputs generated at higher temperatures to enhance decision-making, a property that is not as pronounced in the Solo and Group conditions. This robustness is particularly valuable in real-world medical domains with high uncertainty and ambiguity in datasets [15]. Additionally, studies have shown that creative diagnostic approaches can mitigate cognitive biases and improve diagnostic accuracy [60], while fostering flexibility and adaptability in decision-making [19]. These insights support the enhanced performance observed under higher temperatures in our framework. However, the future work should explore a wider range of temperatures to fully understand the robustness and adaptability of our approach.

4.6 Convergence Trends in Consensus Dynamics

There is clear trend towards consensus among MDAGENTS cross various data modalities (Figure 7). The *text+video* modality demonstrates a rapid convergence, reflecting the agents' efficient processing of combined textual and visual cues. On the other hand, the *text+image* and *text-only* modalities display a more gradual decline in entropy, indicating a progressive narrowing of interpretative diversity among the agents. Despite the differing rates and initial conditions, all modalities exhibit convergence of agent opinions over time. This uniformity in reaching consensus highlights the MDAGENTS' capability to integrate and reconcile information. Please refer to Appendix B for a detailed explanation of the entropy calculation.

5 Conclusion

This paper introduces **MDAGENTS**, a framework designed to enhance the utility of LLMs in complex medical decision-making by dynamically structuring effective collaboration models. To reflect the nuanced consultation aspects in clinical settings, MDAGENTS adaptively assigns LLMs either to roles independently or within groups, depending on the task's complexity. This emulation of real-world medical decision processes has been comprehensively evaluated, with MDAGENTS outperforming previous solo and group methods in **7 out of 10** medical benchmarks. The case study illustrates the practical efficacy and collaborative dynamics of our proposed framework, providing insights into how differing expert opinions are synthesized to reach a more accurate diagnosis. This is evidenced by our agents' ability to converge on the correct diagnosis despite initially divergent perspectives. Ablation studies further elucidate the individual contributions of agents and strategies within the system, revealing the critical components and interactions that drive the framework's success. By harnessing the strength of multi-modal reasoning and fostering a collaborative process among LLM agents, our framework opens up new possibilities for enhancing LLM-assisted medical diagnosis systems, pushing the boundaries of automated clinical reasoning.

References

- [1] Monica Agrawal, Stefan Hegselmann, Hunter Lang, Yoon Kim, and David Sontag. Large language models are few-shot clinical information extractors. *arXiv preprint arXiv:2205.12689*, 2022.
- [2] Zaid Al-Ars, Obinna Agba, Zhuoran Guo, Christiaan Boerkamp, Ziyaad Jaber, and Tareq Jaber. Nlice: Synthetic medical record generation for effective primary healthcare differential diagnosis, 2024.
- [3] Seongsu Bae, Daeun Kyung, Jaehee Ryu, Eunbyeol Cho, Gyubok Lee, Sunjun Kweon, Jungwoo Oh, Lei Ji, Eric Chang, Tackeun Kim, et al. Ehrxqa: A multi-modal question answering dataset for electronic health records with chest x-ray images. *Advances in Neural Information Processing Systems*, 36, 2024.
- [4] Michael L Barnett, Dhruv Boddupalli, Shantanu Nundy, and David W Bates. Comparative accuracy of diagnosis by collective intelligence of multiple physicians vs individual physicians. *JAMA network open*, 2(3):e190096–e190096, 2019.
- [5] Michael L Barnett, Nancy L Keating, Nicholas A Christakis, A James O’malley, and Bruce E Landon. Reasons for choice of referral physician among primary care and specialist physicians. *Journal of general internal medicine*, 27:506–512, 2012.
- [6] Ofir Ben-Assuli, Nanda Kumar, Ofer Arazy, and Itamar Shabtai. The use of analytic hierarchy process for measuring the complexity of medical diagnosis. *Health Informatics Journal*, 26(1):218–232, 2020.
- [7] Justin Bitter, Elizabeth van Veen-Berkx, Hein G Gooszen, and Pierre van Amelsvoort. Multidisciplinary teamwork is an important issue to healthcare professionals. *Team Performance Management: An International Journal*, 19(5/6):263–278, 2013.
- [8] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners, 2020.
- [9] Hanjie Chen, Zhouxiang Fang, Yash Singla, and Mark Dredze. Benchmarking large language models on answering and explaining challenging medical questions, 2024.
- [10] Justin Chih-Yao Chen, Swarnadeep Saha, and Mohit Bansal. Reconcile: Round-table conference improves reasoning via consensus among diverse llms, 2024.
- [11] Weize Chen, Yusheng Su, Jingwei Zuo, Cheng Yang, Chenfei Yuan, Chi-Min Chan, Heyang Yu, Yaxi Lu, Yi-Hsin Hung, Chen Qian, Yujia Qin, Xin Cong, Ruobing Xie, Zhiyuan Liu, Maosong Sun, and Jie Zhou. Agentverse: Facilitating multi-agent collaboration and exploring emergent behaviors, 2023.
- [12] Michael Christ, Florian Grossmann, Daniela Winter, Roland Bingisser, and Elke Platz. Modern triage in the emergency department. *Deutsches Ärzteblatt International*, 107(50):892, 2010.
- [13] Jane Cioffi and Roslyn Markham. Clinical decision-making by midwives: managing case complexity. *Journal of advanced nursing*, 25(2):265–272, 1997.
- [14] Jan Clusmann, Fiona R Kolbinger, Hannah Sophie Muti, Zunamys I Carrero, Jan-Niklas Eckardt, Narmin Ghaffari Laleh, Chiara Maria Lavinia Löffler, Sophie-Caroline Schwarzkopf, Michaela Unger, Gregory P Veldhuizen, et al. The future landscape of large language models in medicine. *Communications medicine*, 3(1):141, 2023.
- [15] Nick Craddock and Laurence Mynors-Wallis. Psychiatric diagnosis: impersonal, imperfect and important. *The British Journal of Psychiatry*, 204(2):93–95, 2014.

- [16] Lesley F Degner and Jeffrey A Sloan. Decision making during serious illness: what role do patients really want to play? *Journal of clinical epidemiology*, 45(9):941–950, 1992.
- [17] Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate, 2023.
- [18] Carolyn Ee, James Lake, Joseph Firth, Fiona Hargraves, M De Manincor, Tanya Meade, Wolfgang Marx, and Jerome Sarris. An integrative collaborative care model for people with mental illness and physical comorbidities. *International Journal of Mental Health Systems*, 14:1–16, 2020.
- [19] Arthur S Elstein and Alan Schwarz. Clinical problem solving and diagnostic decision making: selective review of the cognitive literature. *BMJ*, 324(7339):729–732, 2002.
- [20] Glyn Elwyn, Dominick Frosch, Richard Thomson, Natalie Joseph-Williams, Amy Lloyd, Paul Kinnersley, Emma Cording, Dave Tomson, Carole Dodd, Stephen Rollnick, et al. Shared decision making: a model for clinical practice. *Journal of general internal medicine*, 27:1361–1367, 2012.
- [21] Nancy E Epstein. Multidisciplinary in-hospital teams improve patient outcomes: A review. *Surgical neurology international*, 5(Suppl 7):S295, 2014.
- [22] Zhihao Fan, Jialong Tang, Wei Chen, Siyuan Wang, Zhongyu Wei, Jun Xi, Fei Huang, and Jingren Zhou. Ai hospital: Interactive evaluation and collaboration of llms as intern doctors for clinical diagnosis, 2024.
- [23] Yao Fu, Hao Peng, Tushar Khot, and Mirella Lapata. Improving language model negotiation with self-play and in-context learning from ai feedback, 2023.
- [24] Isaac R Galatzer-Levy, Daniel McDuff, Vivek Natarajan, Alan Karthikesalingam, and Matteo Malgaroli. The capability of large language models to measure psychiatric functioning. *arXiv preprint arXiv:2308.01834*, 2023.
- [25] Amy L Garcia. Variability in acuity in acute care. *JONA: The Journal of Nursing Administration*, 47(10):476–483, 2017.
- [26] Nicki Gilboy, Paula Tanabe, Debbie Travers, Alexander M Rosenau, et al. Emergency severity index (esi): A triage tool for emergency department care, version 4. implementation handbook 2012 edition. *AHRQ publication*, 12, 2011.
- [27] David Grembowski, Judith Schaefer, Karin E Johnson, Henry Fischer, Susan L Moore, Ming Tai-Seale, Richard Ricciardi, James R Fraser, Donald Miller, Lisa LeRoy, et al. A conceptual model of the role of complexity in the care of patients with multiple chronic conditions. *Medical care*, 52:S7–S14, 2014.
- [28] Yu Gu, Robert Tinn, Hao Cheng, Michael Lucas, Naoto Usuyama, Xiaodong Liu, Tristan Naumann, Jianfeng Gao, and Hoifung Poon. Domain-specific language model pretraining for biomedical natural language processing. *ACM Transactions on Computing for Healthcare*, 3(1):1–23, October 2021.
- [29] Deepak Gupta, Kush Attal, and Dina Demner-Fushman. A dataset for medical instructional video classification and question answering, 2022.
- [30] Xuehai He, Yichen Zhang, Luntian Mou, Eric Xing, and Pengtao Xie. Pathvqa: 30000+ questions for medical visual question answering, 2020.
- [31] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.
- [32] Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, Lingfeng Xiao, Chenglin Wu, and Jürgen Schmidhuber. Metagpt: Meta programming for a multi-agent collaborative framework, 2023.

- [33] M. Jimenez-Lara. Reaping the benefits of integrated health care. stanford social innovation review. *Stanford Social Innovation Review*, 2016.
- [34] Di Jin, Eileen Pan, Nassim Oufattolle, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. What disease does this patient have? a large-scale open domain question answering dataset from medical exams, 2020.
- [35] Qiao Jin, Bhuwan Dhingra, Zhengping Liu, William Cohen, and Xinghua Lu. Pubmedqa: A dataset for biomedical research question answering. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2567–2577, 2019.
- [36] Qiao Jin, Yifan Yang, Qingyu Chen, and Zhiyong Lu. Genegpt: augmenting large language models with domain tools for improved access to biomedical information. *Bioinformatics*, 40(2), February 2024.
- [37] Emre Kıcıman, Robert Ness, Amit Sharma, and Chenhao Tan. Causal reasoning and large language models: Opening a new frontier for causality. *arXiv preprint arXiv:2305.00050*, 2023.
- [38] Yubin Kim, Xuhai Xu, Daniel McDuff, Cynthia Breazeal, and Hae Won Park. Health-llm: Large language models for health prediction via wearable sensor data, 2024.
- [39] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *Advances in neural information processing systems*, 35:22199–22213, 2022.
- [40] Tiffany H. Kung, Morgan Cheatham, Arielle Medenilla, Czarina Sillos, Lorie De Leon, Camille Elepaño, Maria Madriaga, Rimel Aggabao, Giezel Diaz-Candido, James Maningo, and Victor Tseng. Performance of chatgpt on usmle: Potential for ai-assisted medical education using large language models. *PLOS Digital Health*, 2(2):1–12, 02 2023.
- [41] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, et al. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466, 2019.
- [42] Danielle L LaFrance, Mary Jane Weiss, Ellie Kazemi, Joanne Gerenser, and Jacqueline Dobres. Multidisciplinary teaming: Enhancing collaboration through increased understanding. *Behavior analysis in practice*, 12(3):709–726, 2019.
- [43] Beibin Li, Konstantina Mellou, Bo Zhang, Jeevan Pathuri, and Ishai Menache. Large language models for supply chain optimization. *arXiv preprint arXiv:2307.03875*, 2023.
- [44] Guohao Li, Hasan Abed Al Kader Hammoud, Hani Itani, Dmitrii Khizbulin, and Bernard Ghanem. Camel: Communicative agents for "mind" exploration of large language model society, 2023.
- [45] Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. Encouraging divergent thinking in large language models through multi-agent debate, 2023.
- [46] Zijun Liu, Yanzhe Zhang, Peng Li, Yang Liu, and Diyi Yang. Dynamic llm-agent network: An llm-agent collaboration framework with agent team optimization, 2023.
- [47] Valentin Liévin, Christoffer Egeberg Hother, Andreas Geert Motzfeldt, and Ole Winther. Can large language models reason about medical questions?, 2023.
- [48] Daniel McDuff, Mike Schaeckermann, Tao Tu, Anil Palepu, Amy Wang, Jake Garrison, Karan Singhal, Yash Sharma, Shekoofeh Azizi, Kavita Kulkarni, et al. Towards accurate differential diagnosis with large language models. *arXiv preprint arXiv:2312.00164*, 2023.
- [49] Ateev Mehrotra, Christopher B Forrest, and Caroline Y Lin. Dropping the baton: specialty referrals in the united states. *The Milbank Quarterly*, 89(1):39–68, 2011.

- [50] Michael Moor, Oishi Banerjee, Zahra Shakeri Hossein Abad, Harlan M. Krumholz, Jure Leskovec, Eric J. Topol, and Pranav Rajpurkar. Foundation models for generalist medical artificial intelligence. *Nature*, 616(7956):259–265, 2023.
- [51] Mark A Musen, Blackford Middleton, and Robert A Greenes. Clinical decision-support systems. In *Biomedical informatics: computer applications in health care and biomedicine*, pages 795–840. Springer, 2021.
- [52] Harsha Nori, Nicholas King, Scott Mayer McKinney, Dean Carignan, and Eric Horvitz. Capabilities of gpt-4 on medical challenge problems, 2023.
- [53] Harsha Nori, Yin Tat Lee, Sheng Zhang, Dean Carignan, Richard Edgar, Nicolo Fusi, Nicholas King, Jonathan Larson, Yuanzhi Li, Weishung Liu, Renqian Luo, Scott Mayer McKinney, Robert Osazuwa Ness, Hoifung Poon, Tao Qin, Naoto Usuyama, Chris White, and Eric Horvitz. Can generalist foundation models outcompete special-purpose tuning? case study in medicine. November 2023.
- [54] OpenAI. Gpt-4 technical report, 2024.
- [55] Anand K Parekh, Richard A Goodman, Catherine Gordon, Howard K Koh, and HHS Interagency Workgroup on Multiple Chronic Conditions. Managing multiple chronic conditions: a strategic framework for improving health outcomes and quality of life. *Public health reports*, 126(4):460–471, 2011.
- [56] Chanwoo Park, Xiangyu Liu, Asuman Ozdaglar, and Kaiqing Zhang. Do llm agents have regret? a case study in online learning and games. *arXiv preprint arXiv:2403.16843*, 2024.
- [57] Khaled Saab, Tao Tu, Wei-Hung Weng, Ryutaro Tanno, David Stutz, Ellery Wulczyn, Fan Zhang, Tim Strother, Chunjong Park, Elahe Vedadi, et al. Capabilities of gemini models in medicine. *arXiv preprint arXiv:2404.18416*, 2024.
- [58] Wenqi Shi, Ran Xu, Yuchen Zhuang, Yue Yu, Jieyu Zhang, Hang Wu, Yuanda Zhu, Joyce Ho, Carl Yang, and May D. Wang. Ehragent: Code empowers large language models for few-shot complex tabular reasoning on electronic health records, 2024.
- [59] Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. Alfworld: Aligning text and embodied environments for interactive learning. *arXiv preprint arXiv:2010.03768*, 2020.
- [60] Hardeep Singh, Denise M Connor, and Gurpreet Dhaliwal. Five strategies for clinicians to advance diagnostic excellence. *BMJ*, 376, 2022.
- [61] Karan Singhal, Shekoofeh Azizi, Tao Tu, S. Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, Ajay Tanwani, Heather Cole-Lewis, Stephen Pfohl, Perry Payne, Martin Seneviratne, Paul Gamble, Chris Kelly, Abubakr Babiker, Nathanael Schärli, Aakanksha Chowdhery, Philip Mansfield, Dina Demner-Fushman, Blaise Agüera y Arcas, Dale Webster, Greg S. Corrado, Yossi Matias, Katherine Chou, Juraj Gottweis, Nenad Tomasev, Yun Liu, Alvin Rajkomar, Joelle Barral, Christopher Semturs, Alan Karthikesalingam, and Vivek Natarajan. Large language models encode clinical knowledge. *Nature*, 620(7972):172–180, 2023.
- [62] Harold C Sox, Michael C Higgins, Douglas K Owens, and Gillian Sanders Schmidler. *Medical decision making*. John Wiley & Sons, 2024.
- [63] J Stairmand, Louise Signal, D Sarfati, C Jackson, L Batten, M Holdaway, and C Cunningham. Consideration of comorbidity in treatment decision making in multidisciplinary cancer team meetings: a systematic review. *Annals of Oncology*, 26(7):1325–1332, 2015.
- [64] Mirac Suzgun and Adam Tauman Kalai. Meta-prompting: Enhancing language models with task-agnostic scaffolding, 2024.
- [65] Miren Taberna, Francisco Gil Moncayo, Enric Jané-Salas, Maite Antonio, Lorena Arribas, Esther Vilajosana, Elisabet Peralvez Torres, and Ricard Mesía. The multidisciplinary team (mdt) approach and quality of care. *Frontiers in oncology*, 10:85, 2020.

- [66] Xiangru Tang, Anni Zou, Zhuosheng Zhang, Ziming Li, Yilun Zhao, Xingyao Zhang, Arman Cohan, and Mark Gerstein. Medagents: Large language models as collaborators for zero-shot medical reasoning, 2024.
- [67] Arsene Fansi Tchango, Rishab Goel, Zhi Wen, Julien Martel, and Joumana Ghosn. Ddxplus: A new dataset for automatic medical diagnosis, 2022.
- [68] Arun James Thirunavukarasu, Darren Shu Jeng Ting, Kabilan Elangovan, Laura Gutierrez, Ting Fang Tan, and Daniel Shu Wei Ting. Large language models in medicine. *Nature medicine*, 29(8):1930–1940, 2023.
- [69] Arun James Thirunavukarasu, Darren Shu Jeng Ting, Kabilan Elangovan, Laura Gutierrez, Ting Fang Tan, and Daniel Shu Wei Ting. Large language models in medicine. *Nature Medicine*, 29(8):1930–1940, 2023.
- [70] Daniel Truhn, Jan-Niklas Eckardt, Dyke Ferber, and Jakob Nikolas Kather. Large language models and multimodal foundation models for precision oncology. *NPJ Precision Oncology*, 8(1):72, 2024.
- [71] Tao Tu, Anil Palepu, Mike Schaekermann, Khaled Saab, Jan Freyberg, Ryutaro Tanno, Amy Wang, Brenna Li, Mohamed Amin, Nenad Tomasev, Shekoofeh Azizi, Karan Singhal, Yong Cheng, Le Hou, Albert Webson, Kavita Kulkarni, S Sara Mahdavi, Christopher Semturs, Juraj Gottweis, Joelle Barral, Katherine Chou, Greg S Corrado, Yossi Matias, Alan Karthikesalingam, and Vivek Natarajan. Towards conversational diagnostic ai, 2024.
- [72] Sean R Tunis, Daniel B Stryer, and Carolyn M Clancy. Practical clinical trials: increasing the value of clinical research for decision making in clinical and health policy. *Jama*, 290(12):1624–1632, 2003.
- [73] Dave Van Veen, Cara Van Uden, Louis Blankemeier, Jean-Benoit Delbrouck, Asad Aali, Christian Bluethgen, Anuj Pareek, Małgorzata Polacin, Eduardo Pontes Reis, Anna Seehofnerová, Nidhi Rohatgi, Poonam Hosamani, William Collins, Neera Ahuja, Curtis P. Langlotz, Jason Hom, Sergios Gatidis, John Pauly, and Akshay S. Chaudhari. Adapted large language models can outperform medical experts in clinical text summarization. *Nature Medicine*, 30(4):1134–1142, February 2024.
- [74] Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Ji-Rong Wen. A survey on large language model based autonomous agents, 2024.
- [75] Sheng Wang, Zihao Zhao, Xi Ouyang, Qian Wang, and Dinggang Shen. Chatcad: Interactive computer-aided diagnosis on medical image using large language models, 2023.
- [76] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022.
- [77] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models, 2023.
- [78] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- [79] Thomas Weida and Jane Weida. Outpatient e/m coding simplified. *Family Practice Management*, 29(1):26–31, 2022.
- [80] Robin M Weinick, Rachel M Burns, and Ateev Mehrotra. Many emergency department visits could be managed at urgent care centers and retail clinics. *Health affairs*, 29(9):1630–1636, 2010.

- [81] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryen W White, Doug Burger, and Chi Wang. Autogen: Enabling next-gen llm applications via multi-agent conversation, 2023.
- [82] Richard C Wuerz, Leslie W Milne, David R Eitel, Debbie Travers, and Nicki Gilboy. Reliability and validity of a new five-level triage instrument. *Academic emergency medicine*, 7(3):236–242, 2000.
- [83] Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, Rui Zheng, Xiaoran Fan, Xiao Wang, Limao Xiong, Yuhao Zhou, Weiran Wang, Changhao Jiang, Yicheng Zou, Xiangyang Liu, Zhangyue Yin, Shihan Dou, Rongxiang Weng, Wensen Cheng, Qi Zhang, Wenjuan Qin, Yongyan Zheng, Xipeng Qiu, Xuanjing Huang, and Tao Gui. The rise and potential of large language model based agents: A survey, 2023.
- [84] Guangzhi Xiong, Qiao Jin, Zhiyong Lu, and Aidong Zhang. Benchmarking retrieval-augmented generation for medicine. *arXiv preprint arXiv:2402.13178*, 2024.
- [85] Cyril Zakka, Akash Chaurasia, Rohan Shad, Alex R. Dalal, Jennifer L. Kim, Michael Moor, Kevin Alexander, Euan Ashley, Jack Boyd, Kathleen Boyd, Karen Hirsch, Curt Langlotz, Joanna Nelson, and William Hiesinger. Almanac: Retrieval-augmented language models for clinical medicine, 2023.
- [86] Andy Zeng, Maria Attarian, Brian Ichter, Krzysztof Choromanski, Adrian Wong, Stefan Welker, Federico Tombari, Aveek Purohit, Michael Ryoo, Vikas Sindhwani, Johnny Lee, Vincent Vanhoucke, and Pete Florence. Socratic models: Composing zero-shot multimodal reasoning with language, 2022.
- [87] Haodi Zhang, Jiahong Li, Yichi Wang, and Yuanfeng Songi. Integrating automated knowledge extraction with large language models for explainable medical decision-making. In *2023 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*, pages 1710–1717. IEEE, 2023.
- [88] Xiaoman Zhang, Chaoyi Wu, Ziheng Zhao, Weixiong Lin, Ya Zhang, Yanfeng Wang, and Weidi Xie. Pmc-vqa: Visual instruction tuning for medical visual question answering, 2023.
- [89] Hongjian Zhou, Fenglin Liu, Boyang Gu, Xinyu Zou, Jinfa Huang, Jinge Wu, Yiru Li, Sam S. Chen, Peilin Zhou, Junling Liu, Yining Hua, Chengfeng Mao, Chenyu You, Xian Wu, Yefeng Zheng, Lei Clifton, Zheng Li, Jiebo Luo, and David A. Clifton. A survey of large language models in medicine: Progress, application, and challenge, 2024.
- [90] Junbin Zhou and Xiao Xu. The difficulty of medical decision-making: should patients be involved? *Hepatobiliary Surgery and Nutrition*, 12(3):407, 2023.

Limitations and Future Works

Despite the successes of our framework in showing promising performance in medical decision-making tasks, we recognize several limitations that open pathways for future research.

Medical Focused Foundation Models. An essential enhancement would be to incorporate the foundation models and systems specifically trained on medical data, like Med-Gemini [57], AMIE [71], and Med-PaLM 2 [61]. These models excel in generating professional medical terminologies, which can facilitate more effective and accurate communication between multiple agents involved in the decision-making process. By leveraging these specialized models, the agents can interact using a shared, precise medical vocabulary, enhancing the system’s overall performance and reliability. This approach not only ensures more medically accurate content generation but also supports better collaboration and understanding among the agents, which is essential for complex medical decision-making tasks.

Patient-Centered Diagnosis. A primary limitation lies in the fact that our current framework operates within the confines of multi-choice question answering and does not account for patient-physician interactions. Real-world medical diagnostics are deeply rooted in patient-centered approaches, where differential diagnosis is informed by a continuous exchange between the patient’s narrative and the physician’s expertise. To bridge this gap, future iterations of our framework will aim to incorporate a more interactive system that not only assists physicians but also engages directly with patients. By embedding real-world interactions between physicians and patients within the feedback loop, the system can be trained to become more nuanced and patient-centric. This adaptation allows it to respond more effectively to individual patient scenarios, leading to higher quality and more personalized medical decision support.

Comparison with Human Clinicians. In this work, we did not include a direct comparison with real clinicians’ performance and their methods of complexity classification. Evaluating how real clinicians classify the complexity of medical cases and comparing their overall performance with our framework remains an open question. Future research should involve simulations and empirical studies that compare each step of our framework with actual medical practice. This comparison will provide deeper insights into the efficacy of MDAgents and identify areas for further refinement and alignment with real-world medical decision-making processes.

Potential Risks and Mitigations. While our framework shows promise, there are potential risks associated with medical hallucinations and the generation of false knowledge. To mitigate these issues in future work, incorporating self-correction processes and a minimum doctor-in-the-loop framework could be beneficial. These measures would help continuously correct the model and improve its accuracy. Additionally, providing confidence scores and uncertainty estimates along with the model’s recommendations would be crucial. This approach would help ensure that the final decision-maker can consider a range of possible options, enhancing the reliability and safety of using MDAgents as a diagnostic support tool.

A Dataset Information

We evaluate multi-agent collaboration frameworks across seven common medical question-answering datasets, which vary in question complexity. Generally, questions are deemed more complex if they involve multiple modalities or entail a lengthy, detailed diagnostic task. Below, we detail each dataset and provide a sample entry:

1. **MedQA.** The MedQA dataset consists of professional medical board exams from the US, Mainland China, and Taiwan [34]. Our study focuses on the English test set, comprising 1,273 questions sourced from the United States Medical Licensing Examination (USMLE). These questions are formatted as multiple-choice text queries with five options. Due to their textual nature and brevity, we categorize these questions as low.
Sample Question: “A 47-year-old female undergoes a thyroidectomy for treatment of Graves’ disease. Post-operatively, she reports a hoarse voice and difficulty speaking. You suspect that this is likely a complication of her recent surgery. What is the embryologic origin of the damaged nerve that is most likely causing this patient’s hoarseness?”

Table 4: Summary of the Datasets. : Text, : Image, : Video. In Appendix A, we provide detailed sample information for each benchmark.

Dataset	Modality	Format	Choice	Testing Size	Domain
MedQA		Question + Answer	A/B/C/D	1273	US Medical Licensing Examination
PubMedQA		Question + Context + Answer	Yes/No/Maybe	500	PubMed paper abstracts
DDxPlus		Question + Answer	A/B/C/D/ . . .	134 K	Pathologies, Symptoms and Antecedents from Patients
SymCat		Question + Answer	A/B/C/D	369 K	Disease-symptom records from public datasets
JAMA		Case + Question + Answer	A/B/C/D	1524	Challenging real-world Clinical Cases from diverse Medical Domains
MedBullets		Case + Question + Answer	A/B/C/D	308	Online Platform Resources for Medical Study
MIMIC-CXR		Question + Answer	Closed Answer	1531	Chest X-ray images and Free-text Reports.
PMC-VQA		Question + Answer	A/B/C/D	50 K	VQA pairs across Images, spanning diverse Modalities and Diseases
Path-VQA		Question + Answer	Yes/No	3391	Open-ended Questions from Pathology Images
MedVidQA		Question + Answer	A/B/C/D	155	First Aids, Medical Emergency, and Medical Education Questions

Options: A: *1st pharyngeal arch*, B: *2nd pharyngeal arch*, C: *3rd pharyngeal arch*, D: *4th pharyngeal arch*, E: *6th pharyngeal arch*

2. **PubMedQA.** PubMedQA is a QA dataset based on biomedical research [35]. It requires yes/no/maybe answers to questions grounded in PubMed abstract. The dataset comprises entries each containing a question, a relevant abstract minus the conclusion, and a ground truth label. We used 50 samples for testing. Given its binary choice format, we consider the complexity of this dataset to be low.

Sample Question: “*Can predilatation in transcatheter aortic valve implantation be omitted?*”

Context: “*The use of a balloon expandable stent valve includes balloon predilatation of the aortic stenosis before valve deployment. The aim of the study was to see whether or not balloon predilatation is necessary in transcatheter aortic valve replacement (TAVI). Sixty consecutive TAVI patients were randomized to the standard procedure or to a protocol where balloon predilatation was omitted. There were no significant differences between the groups regarding early hemodynamic results or complication rates.*”

3. **DDXPlus.** DDXPlus is a medical diagnosis dataset using synthetic patient information and symptoms [67]. Each instance represents a patient, with attributes including age, sex, initial evidences, evidence, multiple options of possible pathologies, and a ground truth diagnosis. Due to its text-only and multiple-choice nature, we consider the complexity to be low.

Sample Patient Information: Age: 96, Sex: F

Evidences: [‘e66’, ‘insp_siffla’, ‘j45’, ‘posttus_emesis’, ‘trav1_@_N’, ‘vaccination’]

Initial evidence: ‘posttus_emesis’

Options: (A) Bronchite (B) Coqueluche

4. **PMC-VQA.** PMC-VQA is a large-scale medical visual question-answering dataset that contains 227K Visual Question Answering (VQA) pairs of 149K images [88]. It is structured as a multiple-choice QA task with one image input accompanying each question. Since the query requires a model to consider both text and image inputs, while maintaining medical expertise, we consider the complexity to be moderate.

Sample Question: *What is the appearance of the hyperintense foci in the basal ganglia on T1-weighted MRI image?*

Image: *PMC8415802 FIG1.jpg*

Options: A: Hypodense, B: Hyperdense, C: Isointense, D: Hypointense

5. **Path-VQA.** PathVQA is a VQA dataset specifically on pathology images [30]. Different from PMC-VQA which consists of multiple choice questions, Path-VQA includes open-ended questions and binary “yes/no” questions. For the purpose of maintaining a standardized accuracy evaluation, we use only the yes/no questions. Similar to PMC-VQA, we consider the complexity to be moderate.

Sample Question: *Was a gravid uterus removed for postpartum bleeding?*

Image: *test_0273.jpg*

6. **MedVidQA.** MedVidQA dataset consists of 3,010 health-related questions with visual answers from validated video sources (e.g. medical school, health institutions, etc). We enhanced the dataset by using GPT-4 to generate multiple-choice answers, including one correct ‘golden answer’ and several false options, expanding its use for training and evaluating automated medical question-answering systems.
Sample Question: How to perform corner stretches to treat neck pain?
Video: h5MvX50zTLM.mp4
Options: A: By bending your knees and touching your toes, B: By performing jumping jacks, C: By leaning into a corner with your elbows up at shoulder level, D: By doing push-ups
7. **SymCat.** SymCat is a synthetic dataset which includes 5 million symptom-condition samples, covering 801 distinct conditions each with 376 potential symptoms dataset [2].
Sample Patient Information: Age: 42, Gender: F, Race: White, Ethnicity: Nonhispanic
Sample Patient Symptoms: Pain:cramping:Abdomen:Lower abdomen::::Worsening after meals:::32; Altered stool:Lumpy::::::Often:44; Flatulence:::::::30; Bloating:::::::41
Options: "A": "Asthma", "B": "IBS (Constipation type)", "C": "Viral meningitis (Varicella zoster virus)", "D": "Bacterial (Gastro)enteritis (Yersinia infection most likely)"
8. **JAMA.** JAMA includes 1524 clinical cases collected from the JAMA Network Clinical Challenge archive, which are summaries of actual challenging clinical cases. Each sample is framed as a question, with a long case description and four options [9].
Sample Case: A 62-year-old woman undergoing peritoneal dialysis (PD) for kidney failure due to IgA nephropathy presented to the PD clinic with a 1-day history of severe abdominal pain and cloudy PD fluid. Seven days prior, she inadvertently broke aseptic technique when tightening a leaking connection of her PD catheter tubing. On presentation, she was afebrile and had normal vital signs. Physical examination revealed diffuse abdominal tenderness. Cloudy fluid that was drained from her PD catheter was sent for laboratory analysis (Table 1). Await peritoneal dialysis fluid culture results before starting intraperitoneal antibiotics.
Sample Question: What Would You Do Next?
"A": "Administer empirical broad-spectrum intraperitoneal antibiotics", "B": "Administer empirical broad-spectrum intravenous antibiotics", "C": "Await peritoneal dialysis fluid culture results before starting intraperitoneal antibiotics", "D": "Send blood cultures"
9. **Medbullets.** Medbullets comprises 308 USMLE Step 2/3 style questions collected from open-access tweets on X (formerly Twitter) since April 2022. The difficulty is comparable to that of Step 2/3 exams, which emulate common clinical scenarios [9].
Sample Question: A 2-week-old boy is evaluated by his pediatrician for abnormal feet. The patient was born at 39 weeks via vaginal delivery to a G1P1 29-year-old woman. The patient has been breastfeeding and producing 5 stools/day. He is otherwise healthy. His temperature is 99.500b0F (37.500b0C), blood pressure is 60/38 mmHg, pulse is 150/min, respirations are 24/min, and oxygen saturation is 98% on room air. A cardiopulmonary exam is notable for a benign flow murmur. A musculoskeletal exam reveals the findings shown in Figure A. Which of the following is the most appropriate next step in management?
Options: "A": "Botulinum toxin injections", "B": "Reassurance and reassessment in 1 month", "C": "Serial casting", "D": "Surgical pinning",
10. **MIMIC-CXR-VQA.** MIMIC-CXR is a large-scale visual question-answering dataset of 377,110 chest radiographs. It was obtained from 227,827 imaging studies sourced from the BIDMC between 2011-2016. It includes patient identifiers which can be linked to MIMIC-IV [3].
Sample Question: Are there any abnormalities in the upper mediastinum?
Image: p11/p11218589/s59138139/7a1e4762-c176bd78-6281fe5c-6b0c9734-e9a4c8f1.jpg

B Entropy Calculation for Consensus Dynamics

The entropy H serves as an indicator of consensus progression among the agents. It is quantified as:

$$H = - \sum_{i=1}^M p(x_i) \log_2 p(x_i) \quad (\text{B.1})$$

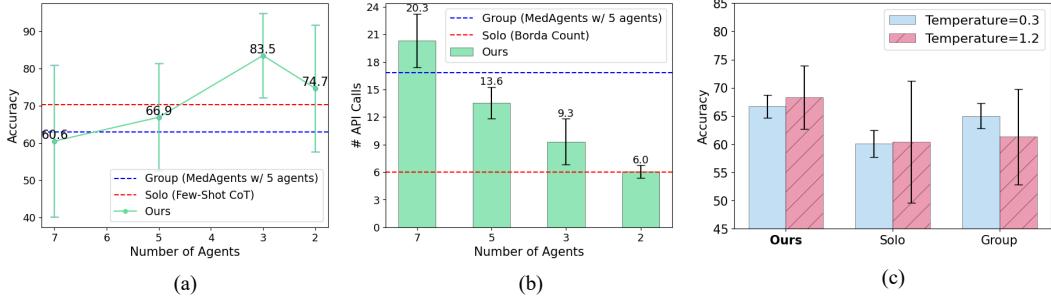


Figure 6: Impact of the number of agents on (a) Accuracy, (b) Number of API Calls on medical benchmarks with GPT-4 (V) and (c) Performance of three different settings under low ($T=0.3$) and high ($T=1.2$) temperatures on medical benchmarks. Our Adaptive setting shows better robustness to different temperatures and even takes advantage of higher temperatures.

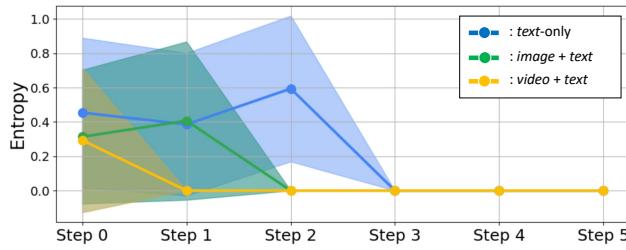


Figure 7: An illustration of consensus entropy in group collaboration process of MDAgents (w/ Gemini-Pro (Vision), $N=30$ for each dataset) on medical benchmarks with different modality inputs.

where M is the total number of unique answers, x_i represents a unique answer, and $p(x_i)$ is the probability of occurring among all answers. This calculation helps to measure the degree of agreement among the agents over time, with lower entropy indicating higher consensus. The trends in entropy across different data modalities provide insights into how quickly and effectively MDAgents can reach a unified decision.

C Prompt Templates

C.1 A single agent setting

```
Few-shot multiple choice questions
{{instruction}}
The following are multiple choice questions (with answers) about medical knowledge.
{{few_shot_examples}}
{{context}} **Question:** {{question}} {{answer_choices}} **Answer:**(
```

```
Chain-of-Thought multiple choice questions
{{instruction}}
The following are multiple choice questions (with answers) about medical knowledge.
{{few_shot_examples w/ CoT Solutions}}
{{context}} **Question:** {{question}} {{answer_choices}} **Answer:**(
```

Ensemble Refinement multiple choice questions

`{{instruction}}`

The following are multiple choice questions (with answers) about medical knowledge.

`{{few_shot_examples w/ CoT Solutions}}`

`{{context}} **Question:** {{question}} {{answer_choices}}`

`{{reasoning_paths}} **Answer:**(`

Medprompt multiple choice questions

`{{instruction}}`

The following are multiple choice questions (with answers) about medical knowledge.

`{{few_shot_examples w/ CoT Solutions from similarity calculation}}`

for N times do

`{{context}} **Question:** {{question}} {{shuffled_answer_choices}}`

`**Answer:**(`

C.2 Multi-agent setting

Complexity check prompt

You are a medical expert who conducts initial assessment and your job is to decide the difficulty/complexity of the medical query.

Now, given the medical query as below, you need to decide the difficulty/complexity of it:

`{{question}}`

Please indicate the difficulty/complexity of the medical query among below options:

1) low: a PCP or general physician can answer this simple medical knowledge checking question without relying heavily on consulting other specialists.

2) moderate: a PCP or general physician can answer this question in consultation with other specialist in a team.

3) high: Team of multi-departmental specialists can answer to the question which requires specialists consulting to another department (requires a lot of team effort to treat the case).

`**Answer:**(`

Recruiter prompt

You are an experienced medical expert who recruits a group of experts with diverse identity and ask them to discuss and solve the given medical query.

Now, given the medical query as below, you need to decide the difficulty/complexity of it:

`{{question}}`

You can recruit up to N experts in different medical expertise. Considering the medical question and the options for the answer, what kind of experts will you recruit to better make an accurate answer?

Also, you need to specify the communication structure between experts (e.g., Pulmonologist == Neonatologist == Medical Geneticist == Pediatrician > Cardiologist)

For example, if you want to recruit five experts, you answer can be like:

`{{exemplars}}`

Please answer in above format, and do not include your reason.

`**Answer:**(`

Agent initialization prompt

You are a {{role}} who {{description}}. Your job is to collaborate with other medical experts in a team.

Agent interaction prompt

Given the opinions from other medical agents in your team, please indicate whether you want to talk to any expert (yes/no). If not, provide your opinion. {{opinions}}

Next, indicate the agent you want to talk to: {{agent_list}}

Remind your medical expertise and leave your opinion to an expert you chose. Deliver your opinion once you are confident enough and in a way to convince other expert with a short reason.

Final decision prompt

You are a final medical decision maker who reviews all opinions from different medical experts and makes the final decision.

Given the {{inputs}}, please review the {{inputs}} and make the final answer to the question by {{decision_methods}} (e.g., majority voting, ensemble refinement).

Answer:(

Algorithm 1 Adaptive Medical Decision-making Framework

Require: Problem Q

```
1:  $Complexity \leftarrow \text{COMPLEXITYCHECK}(Q)$                                  $\triangleright$  Determine the complexity of the medical query
2: if  $Complexity = low$  then
3:    $Agent \leftarrow \text{RECRUIT}(Q, Complexity)$                                  $\triangleright$  Recruit a Primary Care Clinician agent
4:    $ans \leftarrow Agent(Q)$ 
5: else if  $Complexity = moderate$  then
6:    $MDT \leftarrow \text{RECRUIT}(Q, Complexity)$                                  $\triangleright$  Recruit a Multi-disciplinary Team
7:    $Agent \leftarrow \text{RECRUIT}(Q, Complexity, MDT)$ 
8:    $r \leftarrow 0$ 
9:    $Consensus \leftarrow \text{False}$ 
10:   $Interaction \leftarrow []$ 
11:  while  $r \leq R$  and not  $Consensus$  do
12:     $Consensus, Log \leftarrow \text{COLLABORATIVEDISCUSSION}(Q, MDT)$                  $\triangleright$  Iterative discussions
13:    if not  $Consensus$  then
14:      for all  $Agent \in MDT$  do
15:         $Feedback \leftarrow \text{Moderator}(Interaction, Agent)$                                  $\triangleright$  Review and provide feedback
16:         $Agent.\text{UPDATE}(Feedback)$                                                $\triangleright$  Update the feedback
17:      end for
18:       $Interaction \leftarrow Interaction + [Log] + [Feedback]$ 
19:    end if
20:     $r \leftarrow r + 1$ 
21:  end while
22:   $ans \leftarrow Agent(Q, Interaction)$                                           $\triangleright$  Moderator agent makes the final decision
23: else
24:    $ICT \leftarrow \text{RECRUIT}(Q, Complexity)$                                  $\triangleright$  Recruit an Integrated Care Team
25:    $Reports \leftarrow []$ 
26:   for  $Team \in ICT$  do
27:      $Report \leftarrow \text{GENERATEREPORT}(Q, Team)$                                  $\triangleright$  Each Team curates a report
28:      $Reports \leftarrow Reports + [Report]$ 
29:   end for
30:    $ans \leftarrow Agent(Q, Reports)$                                              $\triangleright$  Final decision made
31: end if
32: return  $ans$ 
```

Table 5: Accuracy (%) on Medical benchmarks with **Solo** (⌚) setting. **Bold** represents the best and Underlined represents the second best performance for each benchmark and model.

Medical Knowledge Retrieval Datasets						
Category	Method	MedQA	PubMedQA	Path-VQA	PMC-VQA	MedVidQA
GPT-3.5	Zero-shot	48.5 ± 3.3	<u>56.8</u> ± 12.0	-	-	-
	Few-shot	47.8 ± 16.4	<u>59.0</u> ± 1.0	-	-	-
	+ CoT	54.2 ± 8.9	<u>49.7</u> ± 11.9	-	-	-
	+ CoT-SC	60.5 ± 3.2	<u>49.4</u> ± 13.5	-	-	-
	ER	60.2 ± 4.0	<u>52.9</u> ± 15.9	-	-	-
	Medprompt	60.1 ± 10.8	<u>59.8</u> ± 8.5	-	-	-
GPT-4(V)	Zero-shot	75.0 ± 1.3	<u>61.5</u> ± 2.2	<u>57.9</u> ± 1.6	49.0 ± 3.7	-
	Few-shot	72.9 ± 11.4	<u>63.1</u> ± 11.7	<u>57.5</u> ± 4.5	52.2 ± 2.0	-
	+ CoT [78]	82.5 ± 4.9	<u>57.6</u> ± 9.2	58.6 ± 3.1	51.3 ± 1.5	-
	+ CoT-SC [77]	83.9 ± 2.7	<u>58.7</u> ± 5.0	61.2 ± 2.1	50.5 ± 5.2	-
	ER [61]	81.9 ± 2.1	<u>56.0</u> ± 7.0	61.4 ± 4.1	52.7 ± 2.9	-
	Medprompt [53]	82.4 ± 5.1	51.8 ± 4.6	59.2 ± 5.7	53.4 ± 7.9	-
Gemini-Pro(Vision)	Zero-shot	42.0 ± 10.4	65.2 ± 14.5	45.9 ± 2.8	44.8 ± 2.0	37.9 ± 8.4
	Few-shot	34.0 ± 7.2	<u>55.0</u> ± 0.0	64.5 ± 2.3	48.2 ± 1.0	47.1 ± 8.6
	+ CoT	50.0 ± 6.0	<u>60.2</u> ± 9.0	66.4 ± 11.7	47.1 ± 4.2	48.6 ± 5.5
	+ CoT-SC	52.7 ± 4.6	55.8 ± 8.9	63.6 ± 6.0	46.3 ± 2.8	49.2 ± 8.2
	ER	52.0 ± 7.2	<u>58.4</u> ± 14.2	57.6 ± 8.4	38.4 ± 2.0	48.5 ± 4.1
	Medprompt	45.3 ± 3.1	50.6 ± 5.4	55.0 ± 2.0	41.8 ± 3.0	44.5 ± 2.0
Clinical Reasoning and Diagnostic Datasets						
Category	Method	DDXPlus	SymCat	JAMA	MedBullets	MIMIC-CXR
GPT-3.5	Zero-shot	56.2 ± 4.1	<u>84.0</u> ± 0.0	36.0 ± 3.3	56.0 ± 2.8	-
	Few-shot	48.9 ± 8.5	86.0 ± 2.8	38.0 ± 4.2	59.0 ± 1.4	-
	+ CoT	52.8 ± 5.4	<u>82.0</u> ± 0.0	34.0 ± 2.4	56.0 ± 5.7	-
	+ CoT-SC	37.8 ± 6.1	80.0 ± 2.8	43.0 ± 4.2	63.0 ± 4.2	-
	ER	42.3 ± 6.9	84.0 ± 1.8	44.0 ± 5.7	58.0 ± 0.0	-
	Medprompt	41.2 ± 6.2	86.0 ± 2.6	43.0 ± 1.4	54.0 ± 5.7	-
GPT-4(V)	Zero-shot	70.3 ± 2.0	88.7 ± 2.3	62.0 ± 2.0	67.0 ± 1.4	40.0 ± 5.3
	Few-shot	69.4 ± 1.0	<u>86.7</u> ± 3.1	69.0 ± 4.2	72.0 ± 2.8	35.3 ± 5.0
	+ CoT [78]	72.7 ± 7.7	<u>78.0</u> ± 2.0	66.0 ± 5.7	70.0 ± 0.0	36.2 ± 5.2
	+ CoT-SC [77]	52.1 ± 6.4	83.3 ± 3.1	68.0 ± 2.8	76.0 ± 2.8	51.7 ± 4.0
	ER [61]	61.3 ± 2.4	82.7 ± 2.3	71.0 ± 1.4	76.0 ± 5.7	50.0 ± 0.0
	Medprompt [53]	59.5 ± 17.7	87.3 ± 1.2	70.7 ± 4.3	71.0 ± 1.4	53.4 ± 4.3
Gemini-Pro(Vision)	Zero-shot	49.9 ± 6.5	<u>88.9</u> ± 6.4	42.7 ± 3.9	40.0 ± 1.5	40.0 ± 2.8
	Few-shot	47.1 ± 5.6	89.2 ± 3.4	41.0 ± 1.4	44.0 ± 3.9	39.2 ± 1.2
	+ CoT [78]	65.5 ± 4.9	91.9 ± 3.4	38.0 ± 1.6	52.7 ± 7.1	45.2 ± 6.8
	+ CoT-SC [77]	60.3 ± 2.4	92.0 ± 1.5	46.0 ± 0.0	51.0 ± 4.2	54.9 ± 3.4
	ER [61]	46.7 ± 6.9	58.5 ± 7.5	50.8 ± 5.8	53.2 ± 7.8	53.2 ± 3.5
	Medprompt [53]	58.2 ± 5.5	92.5 ± 4.5	44.4 ± 3.2	54.0 ± 5.7	51.2 ± 1.9

* **CoT**: Chain-of-Thought, **SC**: Self-Consistency, **ER**: Ensemble Refinement

* : text-only, : image+text, : video+text

* All experiments were tested with 3 random seeds

Table 6: Accuracy (%) on Medical benchmarks with **Group** (⌚⌚⌚) setting. **Bold** represents the best performance for each benchmark and model.

Medical Knowledge Retrieval Datasets						
Category	Method	MedQA ⌚⌚⌚	PubMedQA ⌚⌚⌚	Path-VQA ⌚⌚⌚	PMC-VQA ⌚⌚⌚	MedVidQA ⌚⌚⌚
GPT-3.5	Majority Voting	60.4 ± 2.1	68.5 ± 9.6	-	-	-
	Weighted Voting	57.3 ± 3.0	65.8 ± 11.4	-	-	-
	Borda Count	55.3 ± 7.1	70.2 ± 8.8	-	-	-
	MedAgents [66]	56.0 ± 5.3	55.0 ± 1.4	-	-	-
GPT-4(V)	Majority Voting	80.6 ± 2.9	72.2 ± 6.9	56.9 ± 19.7	36.8 ± 6.7	-
	Weighted Voting	78.8 ± 1.1	72.2 ± 6.9	62.1 ± 13.9	25.4 ± 9.0	-
	Borda Count	70.3 ± 8.5	66.9 ± 3.0	61.9 ± 8.1	27.9 ± 5.3	-
	MedAgents [66]	79.1 ± 7.4	69.7 ± 4.7	45.4 ± 8.1	39.6 ± 3.0	-
Gemini-Pro(Vision)	Majority Voting	51.6 ± 2.2	65.3 ± 12.9	58.2 ± 7.3	27.1 ± 5.4	50.8 ± 7.4
	Weighted Voting	52.3 ± 3.3	63.7 ± 10.0	66.4 ± 11.1	20.9 ± 3.8	57.8 ± 2.1
	Borda Count	49.4 ± 9.7	57.7 ± 15.0	68.2 ± 1.8	25.3 ± 8.7	54.5 ± 4.7
	MedAgents [66]	48.4 ± 5.5	63.6 ± 6.0	64.9 ± 12.5	35.1 ± 3.1	61.6 ± 4.8
Multi-agent	Reconcile [10]	81.3 ± 3.0	79.7 ± 3.2	57.5 ± 3.3	31.4 ± 1.2	-
	AutoGen [81]	60.6 ± 5.0	77.3 ± 2.3	43.0 ± 8.9	37.3 ± 6.1	-
	DyLAN [46]	64.2 ± 2.3	73.6 ± 4.2	41.3 ± 1.2	34.0 ± 3.5	-
	Meta-Prompting [64]	80.6 ± 1.2	73.3 ± 2.3	55.3 ± 2.3	42.6 ± 4.2	-
Clinical Reasoning and Diagnostic Datasets						
Category	Method	DDXPlus ⌚⌚⌚	SymCat ⌚⌚⌚	JAMA ⌚⌚⌚	MedBullets ⌚⌚⌚	MIMIC-CXR ⌚⌚⌚
GPT-3.5	Majority Voting	53.6 ± 2.2	83.7 ± 3.3	47.0 ± 1.4	46.0 ± 4.1	-
	Weighted Voting	55.2 ± 2.0	85.9 ± 3.0	49.0 ± 1.4	43.0 ± 8.4	-
	Borda Count	63.9 ± 12.1	84.9 ± 1.6	50.0 ± 1.0	45.0 ± 5.6	-
	MedAgents [66]	47.3 ± 11.0	87.0 ± 4.2	41.0 ± 3.3	56.0 ± 6.2	-
GPT-4(V)	Majority Voting	67.8 ± 4.9	91.9 ± 2.2	70.0 ± 5.7	70.0 ± 0.0	49.5 ± 10.7
	Weighted Voting	65.9 ± 3.3	90.5 ± 2.9	66.1 ± 4.1	66.0 ± 5.7	53.5 ± 2.2
	Borda Count	67.1 ± 6.7	78.0 ± 11.8	61.0 ± 5.6	66.0 ± 2.8	45.3 ± 6.8
	MedAgents [66]	62.8 ± 5.6	90.0 ± 0.0	66.0 ± 5.7	77.0 ± 1.4	43.3 ± 7.0
Gemini-Pro(Vision)	Majority Voting	52.3 ± 15.3	73.5 ± 4.9	47.0 ± 1.4	44.0 ± 4.2	47.9 ± 6.6
	Weighted Voting	54.3 ± 16.9	64.6 ± 6.5	42.0 ± 3.2	43.6 ± 3.9	43.2 ± 2.0
	Borda Count	67.0 ± 27.7	77.3 ± 9.3	37.0 ± 1.8	46.1 ± 3.2	44.6 ± 4.5
	MedAgents [66]	43.0 ± 2.7	80.5 ± 1.9	40.8 ± 2.9	50.5 ± 3.6	38.7 ± 1.5
Multi-agent	Reconcile [10]	68.4 ± 7.4	90.6 ± 2.5	60.7 ± 5.7	59.5 ± 8.7	33.3 ± 3.4
	AutoGen [81]	67.3 ± 11.8	73.3 ± 3.1	64.6 ± 1.2	55.3 ± 3.1	43.3 ± 4.2
	DyLAN [46]	56.4 ± 2.9	75.3 ± 4.6	60.1 ± 3.1	57.3 ± 6.1	38.7 ± 1.2
	Meta-Prompting [64]	52.6 ± 6.1	77.3 ± 2.3	64.7 ± 3.1	49.3 ± 1.2	42.0 ± 4.0

* ⌚⌚⌚: text-only, ⌚⌚⌚: image+text, ⌚⌚⌚: video+text

* All experiments were tested with 3 random seeds

Table 7: Accuracy (%) on Medical benchmarks with **Our** (⌚⌚⌚⌚⌚) method. **Bold** represents the best performance for each benchmark and model.

Medical Knowledge Retrieval Datasets						
Method	MedQA ⌚⌚⌚⌚⌚	PubMedQA ⌚⌚⌚⌚⌚	Path-VQA ⌚⌚⌚⌚⌚	PMC-VQA ⌚⌚⌚⌚⌚	MedVidQA ⌚⌚⌚⌚⌚	
GPT-3.5	64.0 ± 1.6	66.0 ± 5.7	-	-	-	-
GPT-4(V)	88.7 ± 4.0	75.0 ± 1.0	65.3 ± 3.9	56.4 ± 4.5	-	-
Gemini-Pro(Vision)	57.4 ± 1.8	71.0 ± 1.6	72.0 ± 2.3	62.2 ± 7.6	56.2 ± 6.7	-
Clinical Reasoning and Diagnostic Datasets						
Method	DDXPlus ⌚⌚⌚⌚⌚	SymCat ⌚⌚⌚⌚⌚	JAMA ⌚⌚⌚⌚⌚	MedBullets ⌚⌚⌚⌚⌚	MIMIC-CXR ⌚⌚⌚⌚⌚	
GPT-3.5	62.5 ± 6.7	85.7 ± 3.0	48.6 ± 6.5	55.3 ± 4.3	-	-
GPT-4(V)	77.9 ± 2.1	93.1 ± 1.0	70.9 ± 0.3	80.8 ± 1.7	55.9 ± 9.1	-
Gemini-Pro(Vision)	59.2 ± 1.2	65.0 ± 6.1	47.0 ± 2.2	42.9 ± 3.4	48.1 ± 5.8	-

* ⌚⌚⌚⌚⌚: text-only, ⌚⌚⌚⌚⌚: image+text, ⌚⌚⌚⌚⌚: video+text

* All experiments were tested with 3 random seeds

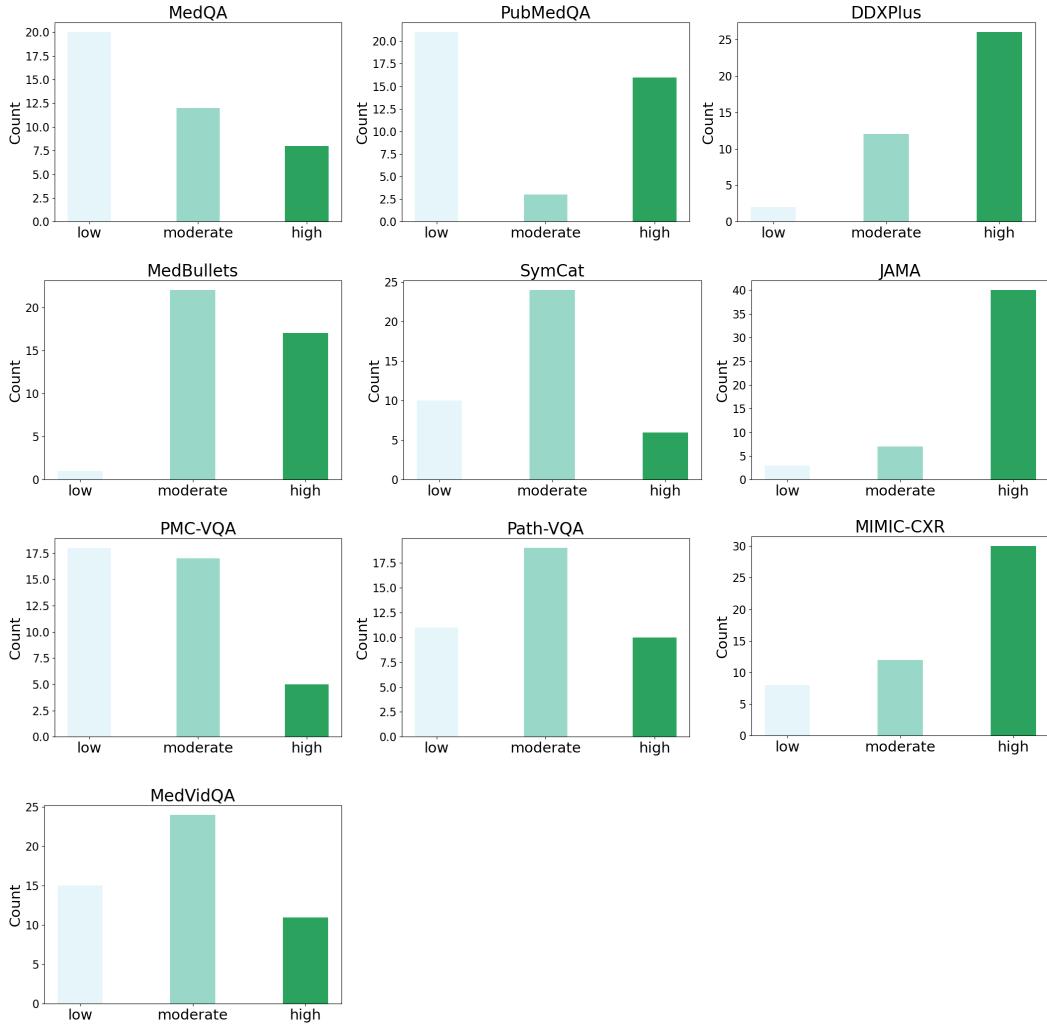


Figure 8: Complexity Distribution for each dataset classified by GPT-4(V) and Gemini-Pro (Vision) (for MedVidQA). The plot illustrates the varying levels of medical complexity across datasets, reflecting the diverse nature of medical question answering, diagnostic reasoning, and medical visual interpretation tasks. For instance, MedQA is categorized under Medical Knowledge Retrieval due to their focus on text-based questions and literature synthesis, while MIMIC-CXR, categorized under Clinical Reasoning and Diagnostic tasks, shows a high complexity distribution due to the need for interpreting detailed radiographic images (See Section in Section 4.1 for the task categorization)

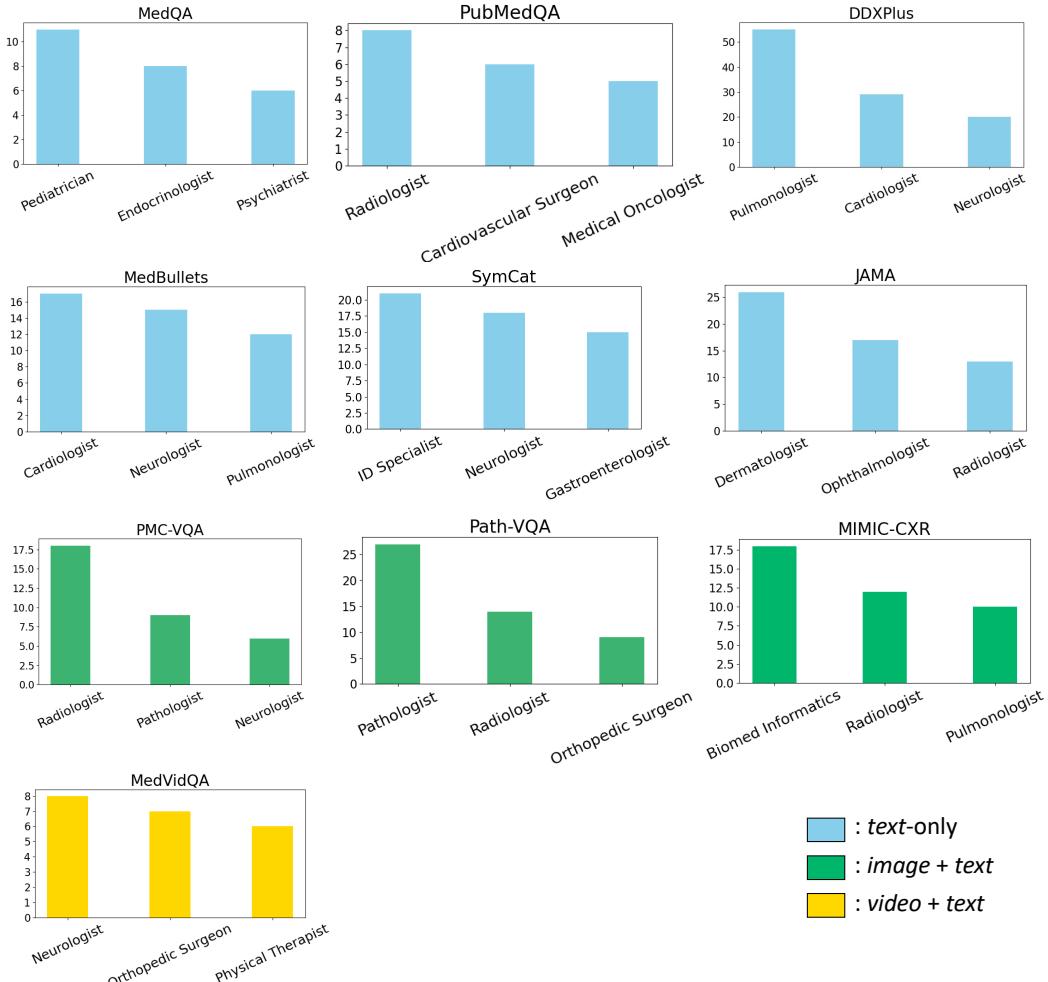


Figure 9: Top-3 most recruited medical experts in each benchmark. The alignment between the dataset characteristics and the recruited experts is evident in several cases. For instance, MIMIC-CXR, which features chest x-ray images, predominantly recruits Radiologists, Pulmonologists, and experts in Biomedical Informatics due to their expertise in interpreting medical imaging.

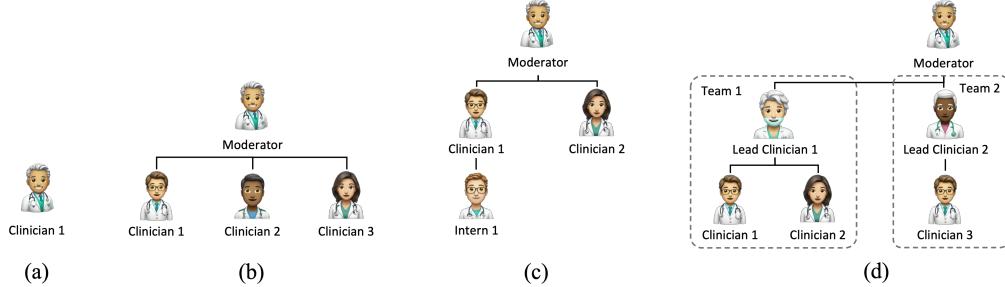


Figure 10: Simplified agent structure examples assigned during the expert recruitment process ranging from (a) A Primary Care Clinician (PCC), (b) Multi-disciplinary Team (MDT), (C) MDT w/ hierarchy to (d) Integrated Care Team (ICT).

D Case Study

D.1 Medical Decision Making Case Studies

MDM requires efforts of both individual expertise and collaboration to navigate the complexities of patient care. Clinicians often face challenging scenarios that necessitate a comprehensive approach, integrating insights from various specialties to arrive at the best possible outcomes.

D.1.1 Real-World Medical Cases

Below are the real-world example cases that could be classified as low, moderate, to high complexity cases.

Case 1: Adjusting Medication Dosage for Chronic Disease (Low Complexity) A 55-year-old female patient with type 2 diabetes visits her PCP for a routine check-up. The patient has been taking 500 mg of metformin orally twice a day and has been adhering to a low-carbohydrate diet. Upon testing with fasting glucose level, the glucose level is above normal. PCP reviewed the current medical dosage and increased the dosage to manage the blood glucose level of the patient.

Case 2: Differential Diagnosis in the Emergency Department (Moderate Complexity) A 40-year-old male patient arrives at the emergency department (ED) with a high fever, severe headache, and muscle pain, raising concerns about a potential infectious disease. The ED physician conducts an initial examination but recognizes the need for a more detailed evaluation to identify the underlying cause. The patient is referred to the infectious disease department for further assessment. An infectious disease specialist, along with the ED physician, reviews the patient's symptoms, travel history, and recent exposures. They collaborate on ordering specific diagnostic tests, including blood cultures and imaging studies. Through this teamed decision-making process, they diagnose the patient with dengue fever and promptly initiate appropriate antiviral treatment.

Case 3: Managing Adverse Responses to Medication in Chronic Disease (High Complexity) A 60-year-old female patient with chronic heart failure has been experiencing new symptoms of shortness of breath and mild fever, suggesting either a complication due to her chronic heart failure or a new infection. The urgent care doctor identifies the severity of the situation and promptly refers the patient to the emergency department of a large hospital, where the patient has triaged to see a cardiologist and an infectious disease doctor for specialized care. The team conducts a detailed review of the patient's medication history and current symptoms, does a physical exam to listen to lung sounds, and orders a few exams including labs, a chest x-ray, echocardiogram, and electrocardiogram. The team identifies that the patient has pulmonary effusion and upper respiratory viral infection.

D.1.2 Medical Cases from MedQA Dataset

Now, let us look at the cases from the MedQA [34] dataset that illustrate either individual PCP or teamed decision-making is crucial in managing medical conditions, ranging from low to high complexity levels of potential cases. These examples highlight the importance of checking the complexity of the case for proper management.

Case 4: Diagnosis by PCP The case below with the “Low Complexity” header is classified as low complexity by a medical doctor. In this case, a PCP can answer this question without consulting a gastroenterologist. The diagnosis of gastric cancer and management based on the manifestation of the disease, that has been described in this question and beyond should be from a gastroenterologist. However, PCPs are expected to have the basic scientific and pathophysiological knowledge that is related to gastric cancer and use that knowledge to solve this problem.

Case 5: Diagnosis and management by single Pediatric Endocrinologist The case below with the “Moderate Complexity” header is classified as moderate complexity by the medical doctor. In this case, a pediatric endocrinologist (specialist) alone can diagnose a patient and have a treatment plan. Note that this patient could have been referred to this pediatric endocrinologist by a PCP who is regularly seeing this patient.

Low Complexity

Question: A 70-year-old man comes to the physician because of a 4-month history of epigastric pain, nausea, and weakness. He has smoked one pack of cigarettes daily for 50 years and drinks one alcoholic beverage daily. He appears emaciated. He is 175 cm (5 ft 9 in) tall and weighs 47 kg (103 lb); BMI is 15 kg/m². He is diagnosed with gastric cancer. Which of the following cytokines is the most likely direct cause of this patient's examination findings?

Answer:

- A) TGF- β
- B) IL-6
- C) IL-2
- D) TNF- β

Moderate Complexity

Question: A 5-year-old girl is brought to the clinic by her mother for excessive hair growth. Her mother reports that for the past 2 months she has noticed hair at the axillary and pubic areas. She denies any family history of precocious puberty and reports that her daughter has been relatively healthy with an uncomplicated birth history. She denies any recent illnesses, weight change, fever, vaginal bleeding, pain, or medication use. Physical examination demonstrates Tanner stage 4 development. A pelvic ultrasound shows an ovarian mass. Laboratory studies demonstrate an elevated level of estrogen. What is the most likely diagnosis?

Answer:

- A) **Granulosa cell tumor**
- B) Idiopathic precocious puberty
- C) McCune-Albright syndrome
- D) Sertoli-Leydig tumor

Case 6: Diagnosis and management by multidisciplinary team The case below is of a “High Complexity” patient primarily having neurological symptoms but with problems with vision, which requires a neurologist to consult to ophthalmology department for further evaluation.

High Complexity

Question: A 63-year-old woman presents to her primary-care doctor for a 2-month history of vision changes, specifically citing the gradual onset of double vision. Her double vision is present all the time and does not get better or worse throughout the day. She has also noticed that she has a hard time keeping her right eye open, and her right eyelid looks ‘droopy’ in the mirror. Physical exam findings during primary gaze are shown in the photo. Her right pupil is 6 mm and poorly reactive to light. The rest of her neurologic exam is unremarkable. Laboratory studies show an Hb A1c of 5.0%. Which of the following is the next best test for this patient?

Answer:

- A) Direct fundoscopy
- B) Intraocular pressures
- C) **MR angiography of the head**
- D) Temporal artery biopsy

D.2 Cases Studies with MDAgents

In this section, we provide two examples of our framework with *moderate* (Figure 11) and *high* (Figure 12) complexity in PMC-VQA (*image+text*) and DDXPlus (*text-only*) respectively. These

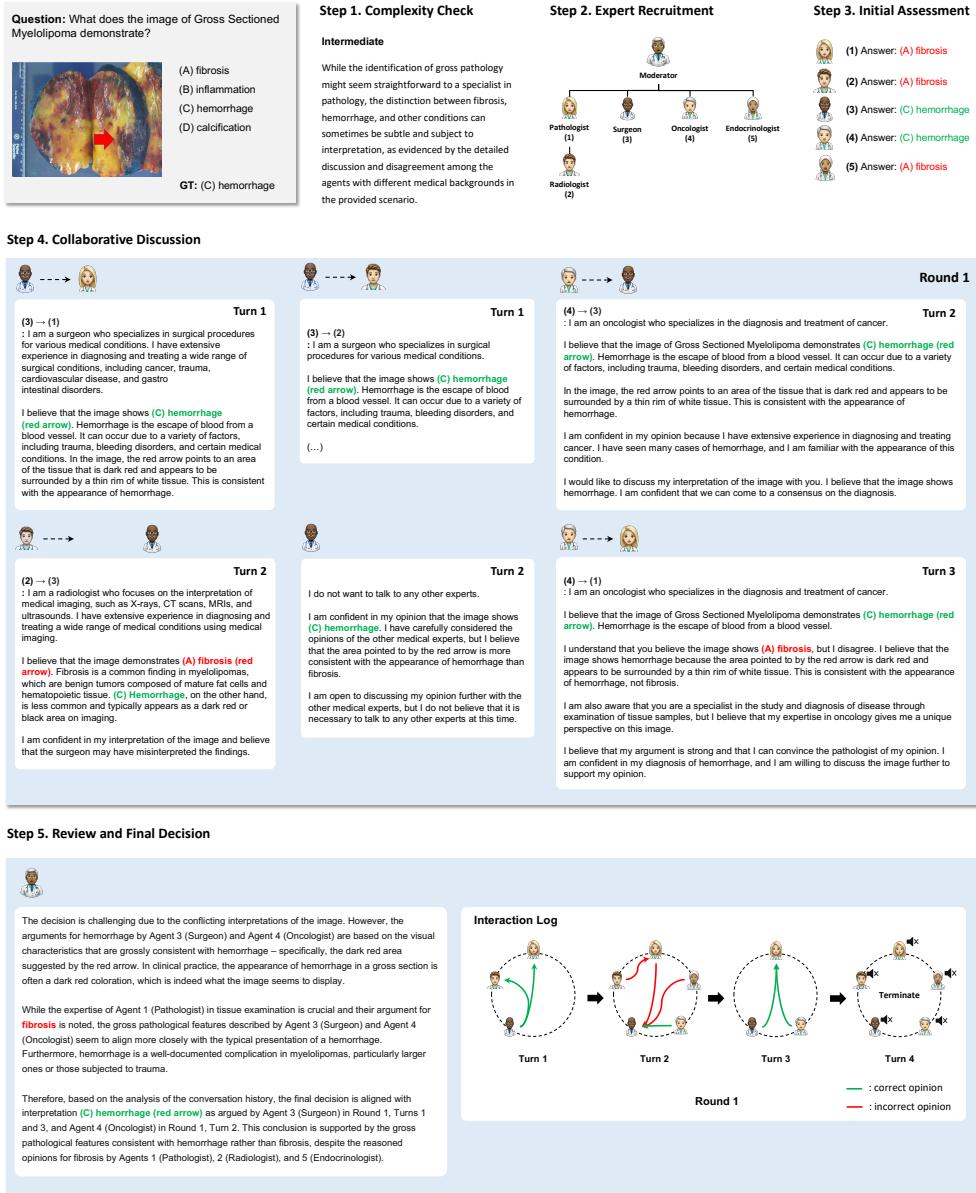


Figure 11: Illustration of our proposed framework in *moderate* complexity setting. Given a medical query (*image + text*) the framework performs reasoning in five steps: (i) complexity check, (ii) expert recruitment, (iii) initial assessment, (iv) collaborative discussion, and (v) review and final decision-making. **Green** text represents the correct answer and the **Red** text represents the incorrect answer.

case studies reveals how our framework provides an environment for agents to collaborate, gather information, moderate and make final decisions in complex medical scenarios.

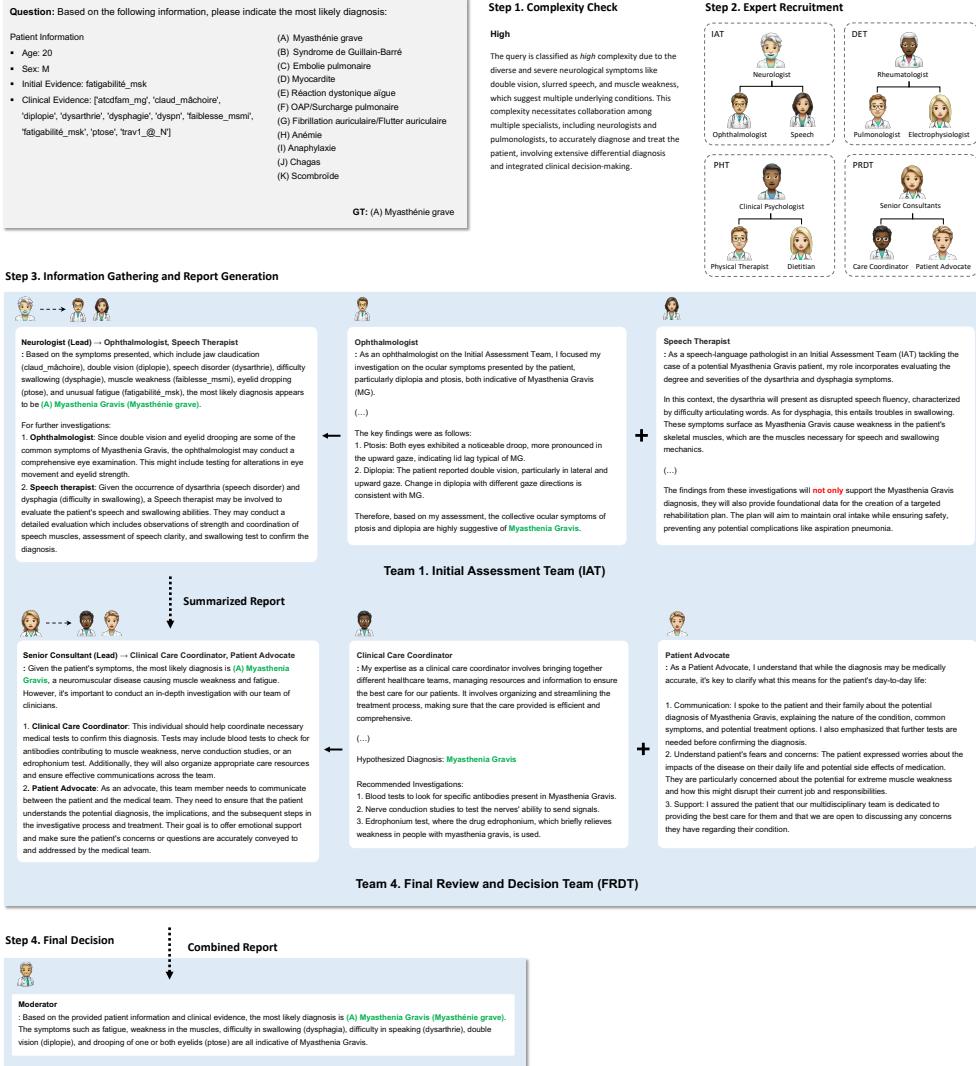


Figure 12: Illustration of our proposed framework in *high* complexity setting. Given a medical query (*text-only*) the framework performs reasoning in four steps: (i) complexity check, (ii) expert recruitment, (iii) information gather and report generation, (iv) final decision. **Green** text represents the correct answer.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state the primary contributions of the paper, including the introduction of the MDAgents framework, its adaptive decision-making structure, and the significant performance improvements demonstrated through experimental results.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The paper includes a section on limitations, discussing the scope of the framework, and areas for future improvement.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: The paper does not include theoretical results that require assumptions or proofs.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper provides comprehensive details on datasets, experimental setups, and methodologies used, ensuring that the results can be reproduced accurately.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The paper includes a link to the code repository and provides sufficient instructions in the supplemental material to reproduce the experiments.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: The paper specifies all relevant experimental details, including data splits, number of samples, and number of seeds, ensuring transparency and reproducibility of the results.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports error bars and includes information on the statistical significance of the experimental results, ensuring a clear understanding of the variability and reliability of the findings.

Guidelines:

- The answer NA means that the paper does not include experiments.

- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The experiments primarily involved inference using API calls to GPT-3.5, GPT-4 (V), and Gemini-Pro (Vision). The type of compute workers, memory, and time of execution are managed by the API providers (OpenAI and Gemini). Details about the number of API calls and the specific configurations used for each experiment are provided, ensuring reproducibility.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research adheres to the NeurIPS Code of Ethics, ensuring responsible conduct throughout the study.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper discusses potential societal impacts, both positive and negative, including the improvement of medical decision-making and the risks associated with misuse or biases in the AI system.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [Yes]

Justification: The paper discusses potential limitations, including the risks of medical hallucinations and false knowledge. It includes guidelines within the main text to ensure responsible use of the MDAgents framework and publicly available datasets, emphasizing ethical usage and adherence to best practices.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The paper properly credits the creators of existing assets used and clearly states the licenses and terms of use.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New Assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: The paper does not introduce any new assets.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and Research with Human Subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: The paper does not involve crowdsourcing or research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: The paper does not involve crowdsourcing or research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.